

Beyond the Basics: Customization Through Configuration






Justin Pihony


DEVELOPER SUPPORT MANAGER @ LIGHTBEND

@JustinPihony









Configuration





Broker


- Scale
- Parallelization
- Fault Tolerance
- Data Retention


Detangling the Configuration Overload





























Configuring the Broker: Core Configuration and Management



How?

kafka-server-start **config/server.properties**

1.1.0+

kafka-configs --bootstrap-server localhost:9092 ...



NAME	DESCRIPTION	TYPE	DEFAULT	VALID VALUES	IMPORTANCE	DYNAMIC UPDATE MODE
------	-------------	------	---------	--------------	------------	------------------------

read-only

Requires server.properties change and broker restart

per-broker

kafka-configs must provide entity-name [BROKER.ID]

cluster-wide

kafka-configs *can* provide entity-default



Minimum Broker Configurations



broker.id

zookeeper.connect

log.dirs



Important Broker Configurations



listeners

format → listener_name://host_name:port


default → PLAINTEXT://:9092




Important Broker Configurations

message.max.bytes \geq **fetch.message.max.bytes**
default \rightarrow 1000012 (1 mb)
topic override = max.message.bytes

BROKER





CONSUMER




Configuring the Broker: Parallelizing via Partitions









`num.partitions = [MULTIPLE OF 6]`







How Many Partitions?







How Many Partitions?



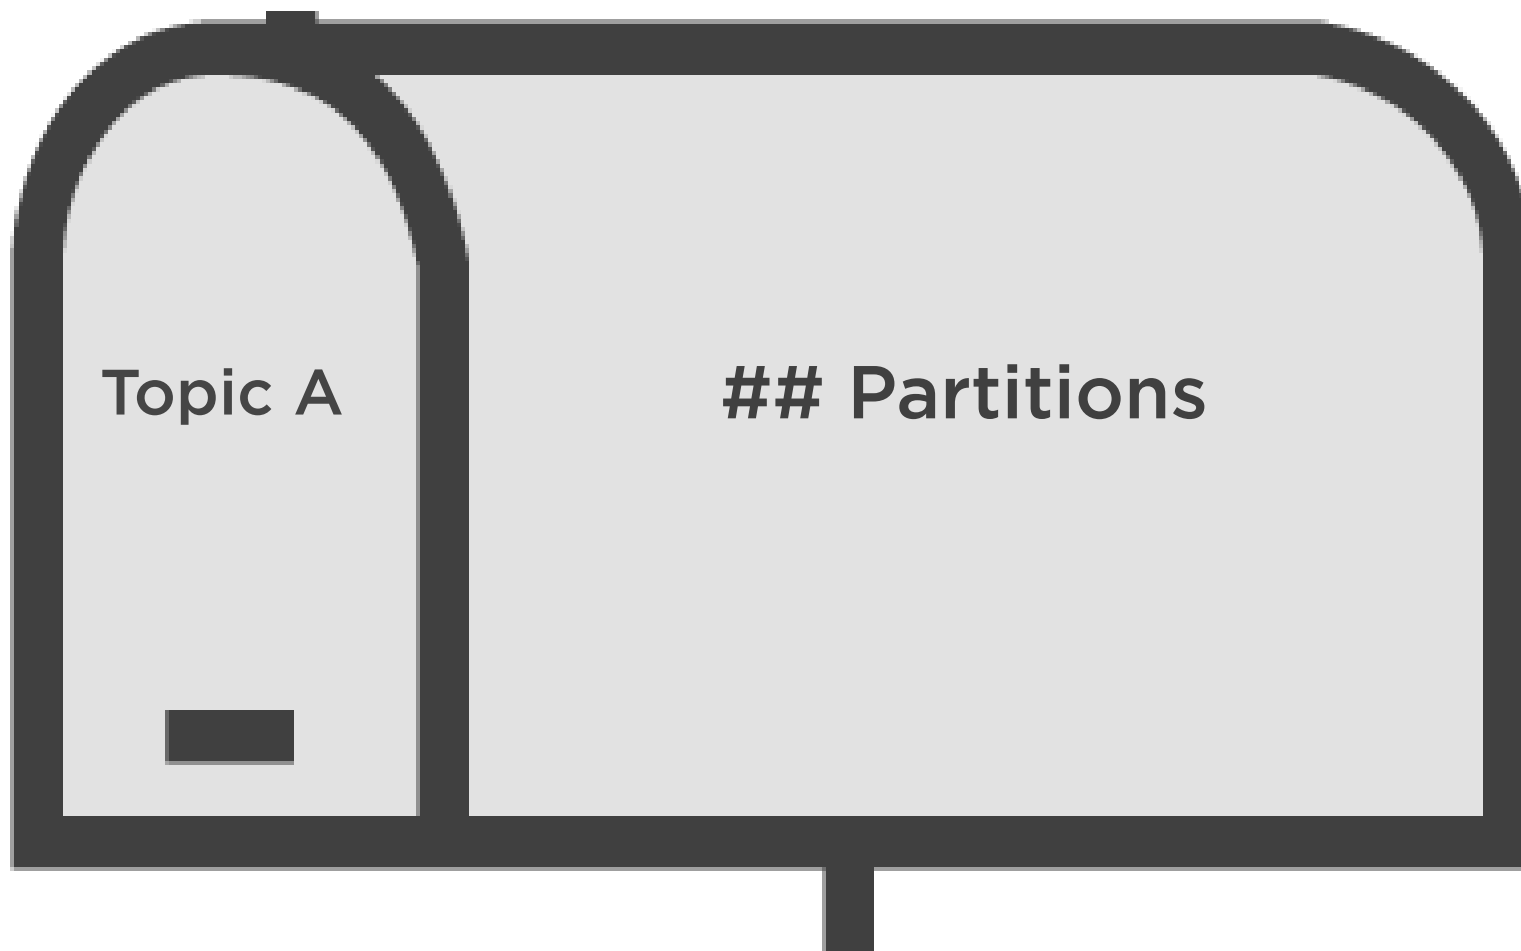


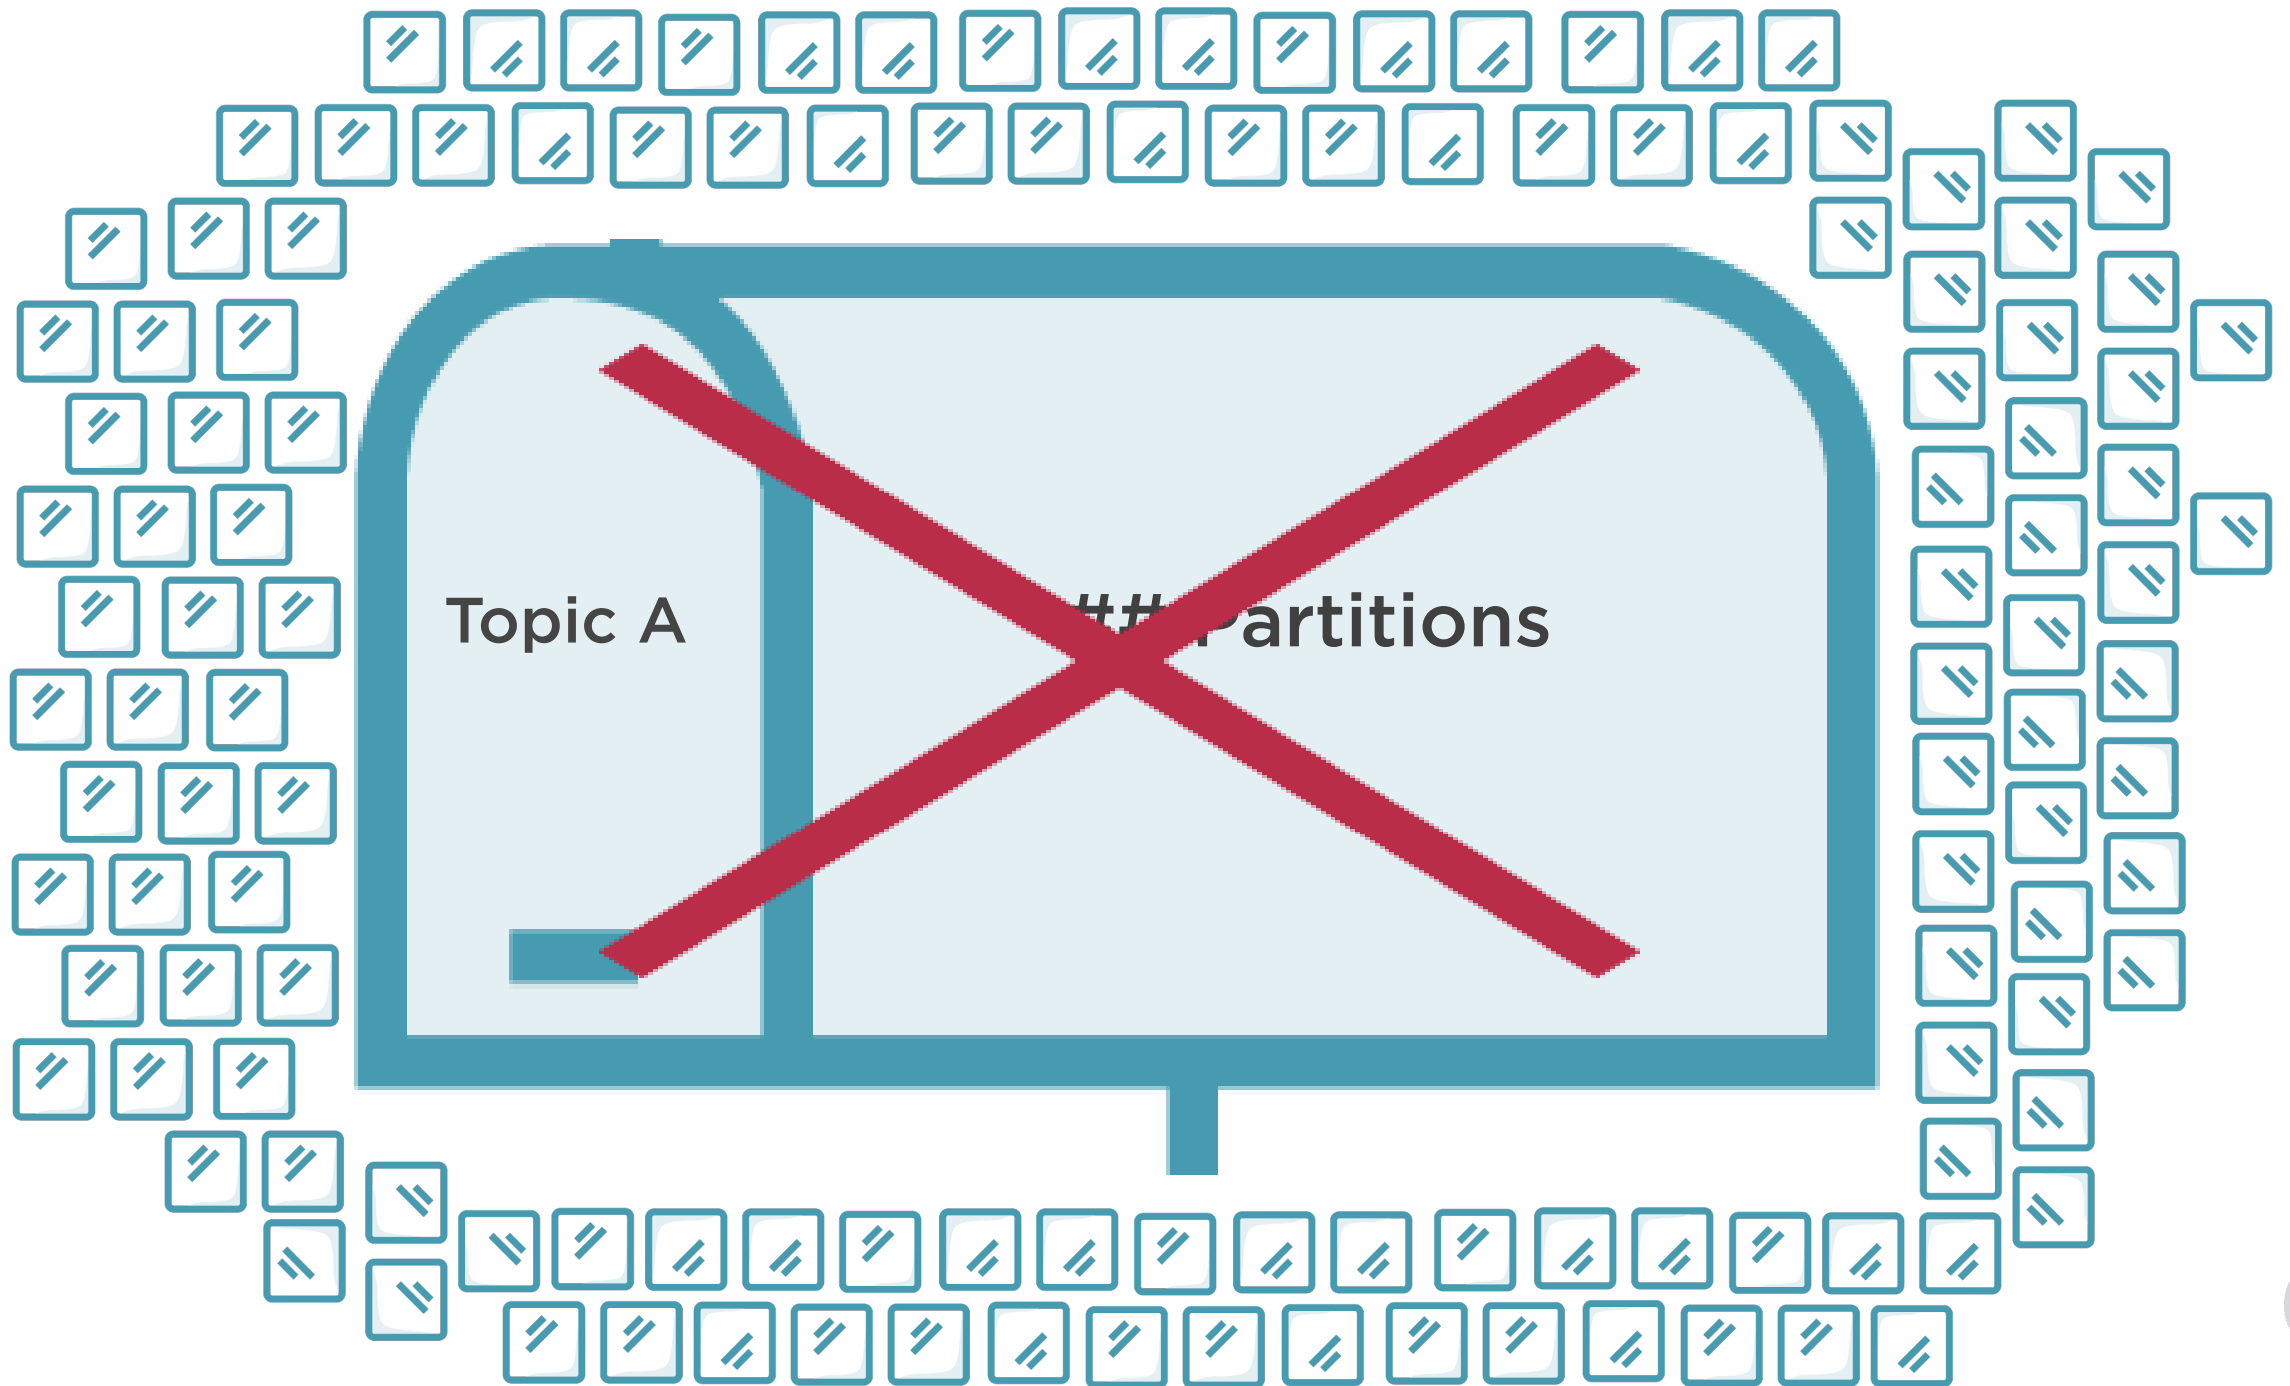














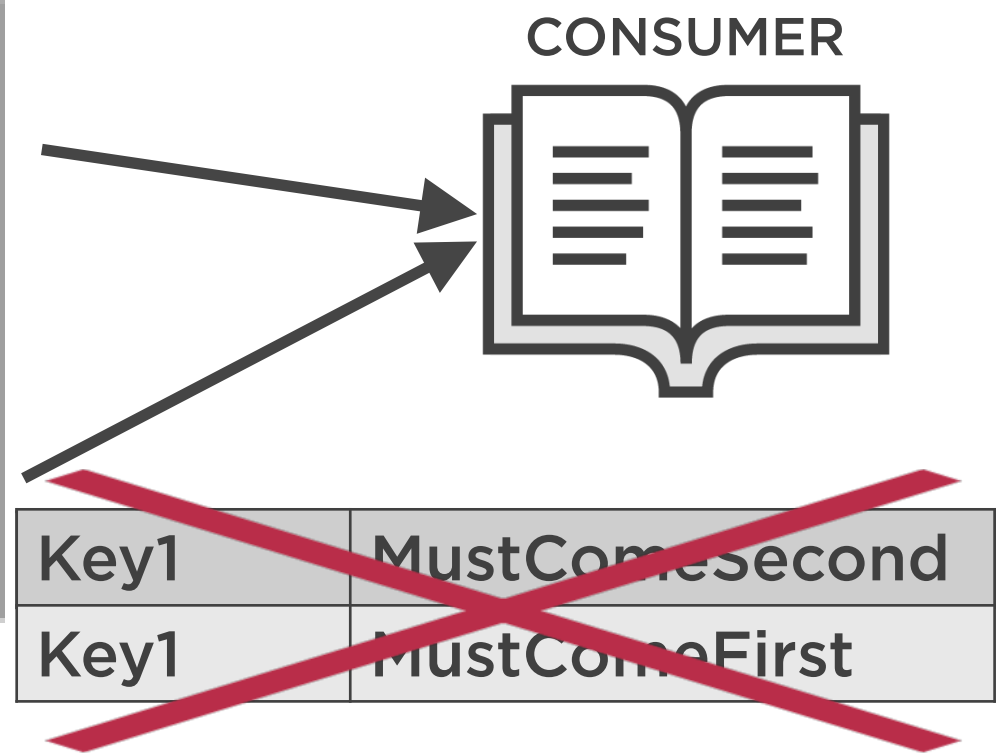


















Broker Config

num.partitions (default = 1)

kafka-topics.sh --create ... **--partitions ##**


auto.create.topics.enable (default = true)



Configuring the Broker: Fault Tolerance via Replication



Replication




Replication



```
kafka-topics.sh --create ... --replication-factor ##
```



Replication




replication-factor = 2




What to Set Replication Factor?




replication-factor




What to Set Replication Factor?




What to Set Replication Factor?



What to Set Replication Factor?



What to Set Replication Factor?




Replication

kafka-topics.sh --create ... **--replication-factor ##**

default.replication.factor (default = 1)
(only for automatic topic creation)




Replication



11010
01011
00101



Replication




`unclean.leader.election.enable`

`>= Kafka 0.11 (default = false)`

`<= Kafka 0.11 (default = true)`


Replication



unclean.leader.election.enable

true


Replication



unclean.leader.election.enable

true


Replication



`unclean.leader.election.enable`

true

Replication




`unclean.leader.election.enable`

`>= Kafka 0.11 (default = false)`


`<= Kafka 0.11 (default = true)`

Replication

Cluster




Replication




`replica.lag.time.max.ms`
10000




Replication




Replication




Replication




Replication



Replication



Replication





Replication

Cluster





`auto.leader.rebalance.enable`


false





Configuring the Broker: Data Retention













Data Management



log.dirs

(server.properties default = tmp/kafka-logs)



5 Partitions



7 Partitions



3 Partitions



0 Partitions



Data Management




log.dirs

(server.properties default = tmp/kafka-logs)



num.recovery.threads.per.data.dir

(default = 1)



Data Management



log.dirs

(server.properties default = tmp/kafka-logs)




num.recovery.threads.per.data.dir

(default = 1)


log.dirs
x
recovery



Partition



Partition



log.segment.bytes (default = 1073741824 (1GB))







`log.retention.bytes`

`log.retention.[hours/minutes/ms]`





log.retention.bytes (per partition)
(default = -1 (unlimited))

log.retention.[hours/minutes/ms]

**# of partions
x
retention**





log.retention.bytes (per partition)

(default = -1 (unlimited))



log.retention.[hours/minutes/ms] (per partition)

(default = 168 hours)




log.retention.bytes (per partition)

(default = -1 (unlimited))


log.retention.[hours/minutes/ms] (per partition)

(default = 168 hours)




last modified time





`log.segment.bytes`






`log.segment.bytes`

`log.segment.ms`
(default = null)





log.segment.bytes

log.segment.ms
(default = null)








`log.segment.bytes` -> 1GB





`log.retention.hours` -> 7 days
`log.segment.bytes` -> 1GB



Resources

- **How to choose the number of topics/partitions**
 - confluent.io/blog/how-choose-number-topics-partitions-kafka-cluster
- **Kafka compaction**
 - kafka.apache.org/documentation/#compaction
- **More on replication**
 - confluent.io/blog/hands-free-kafka-replication-a-lesson-in-operational-simplicity



Summary



Broker

- Scale
- Parallelization
- Fault Tolerance
- Data Retention

