

# Production Hardening: Customizing Towards Reliable Data

---



**Justin Pihony**

DEVELOPER SUPPORT MANAGER @ LIGHTBEND

@JustinPihony



# Hardening & Customizing



## More Configuration

- Producer
- Consumer

## Reliability

## Message Customization




# Configuring the Producer

---




# Producer




# Producer




acks



acks



acks = 0




recordMetadata.offset = -1








acks = 0




acks = 1 (default)




acks = 1 (default)




acks = 1 (default)




acks = 1 (default)




`unclean.leader.election.enable`  
=  
`true`




acks = all (or -1)






retries







retries



retries = 0 (off)




retries = 0 (off)



`delivery.timeout.ms(2.1+)`  
=  
120000  
(default)





```
retries = Integer.MAX_VALUE  
(default in 2.1+)
```



```
delivery.timeout.ms  
=  
120000  
(default)
```



retries = 1



retried = 1

&& total time < ...

delivery.timeout.ms


=

120000

(default)



retries = 1




delivery.timeout.ms  
=  
120000  
(default)

retry.backoff.ms = 100 (default)  
(0.10.2+)



retries = 1



delivery.timeout.ms  
=  
120000  
(default)



retry.backoff.ms = 100 (default)





`retries = 1`

`max.in.flight.requests.per.connection`

`=`

`5`

`(default)`



`retry.backoff.ms = 100 (default)`

`delivery.timeout.ms`

`=`

`120000`


`(default)`






retries = 1

max.in.flight.requests.per.connection  
=  
1



delivery.timeout.ms  
=  
120000  
(default)



retry.backoff.ms = 100 (default)





`retries = 1`

`max.in.flight.requests.per.connection`

`=`

`5`

`(default)`



`retry.backoff.ms = 100 (default)`


`delivery.timeout.ms`

`=`

`120000`

`(default)`







`request.timeout.ms`  
=  
30000  
(default)



# retries


`max.in.flight.requests.per.connection`


$$\frac{\text{request.timeout.ms}}{\text{delivery.timeout.ms}}$$


`retry.backoff.ms`




buffer.memory  
=  
33554432  
(default)



retries

max.in.flight.requests.per.connection




$$\frac{\text{request.timeout.ms}}{\text{delivery.timeout.ms}}$$



retry.backoff.ms




buffer.memory  
=  
33554432  
(default)



retries

max.in.flight.requests.per.connection



$$\frac{\text{request.timeout.ms}}{\text{delivery.timeout.ms}}$$



retry.backoff.ms





`max.block.ms`  
=  
60000  
(default)


---

`buffer.memory`  
=  
33554432  
(default)



retries

`max.in.flight.requests.per.connection`



$$\frac{\text{request.timeout.ms}}{\text{delivery.timeout.ms}}$$



`retry.backoff.ms`



```
block.on.buffer.full  
= (< 0.9)
```

```
true  
(default)
```

```
max.block.ms  
=  
60000  
(default)
```


---

```
buffer.memory  
=  
33554432  
(default)
```



retries

max.in.flight.requests.per.connection


$$\frac{\text{request.timeout.ms}}{\text{delivery.timeout.ms}}$$


retry.backoff.ms





`block.on.buffer.full`

`=`

`false`  
`(default)`

`max.block.ms`

`=`

`60000`  
`(default)`

---

`buffer.memory`


`=`

`33554432`  
`(default)`



retries

`max.in.flight.requests.per.connection`



$$\frac{\text{request.timeout.ms}}{\text{delivery.timeout.ms}}$$




`retry.backoff.ms`



`max.block.ms`  
=  
60000  
(default)


---

`buffer.memory`  
=  
33554432  
(default)



retries

`max.in.flight.requests.per.connection`




$$\frac{\text{request.timeout.ms}}{\text{delivery.timeout.ms}}$$




`retry.backoff.ms`






`batch.size`  
`linger.ms`





~~batch.size~~ | max.request.size  
linger.ms





(broker config)  
message.max.bytes  
>  
batch.size|max.request.size  
linger.ms






retries

`max.in.flight.requests.per.connection`

`max.block.ms`

`buffer.memory`



`ssl/sasl...`




`batch.size` | `max.request.size`  
`linger.ms`

`request.timeout.ms`  

---

`delivery.timeout.ms`



`retry.backoff.ms`



# Configuring the Consumer

---






# Consumer


`poll(TIMEOUT)`




# Consumer



# Consumer





# Consumer




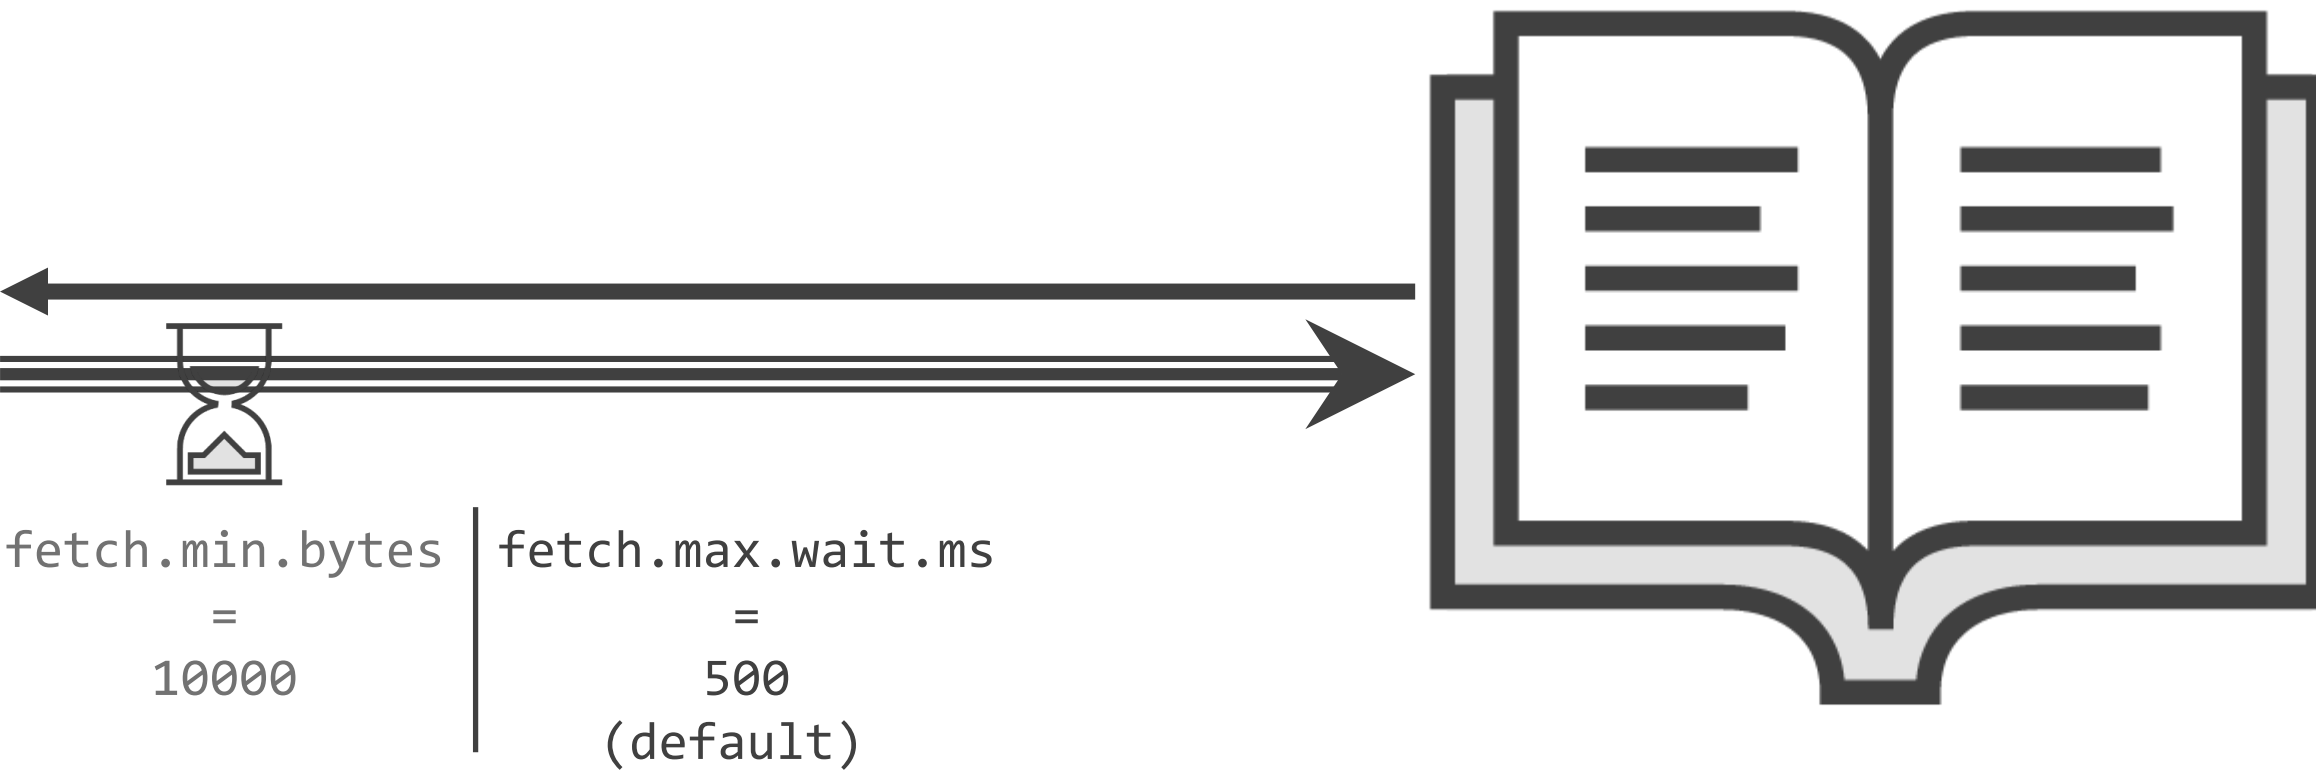
Cluster Load



## Cluster







`fetch.max.bytes` | `max.poll.records`



`fetch.min.bytes` | `fetch.max.wait.ms`





send.buffer.bytes  
=  
102400  
(default)

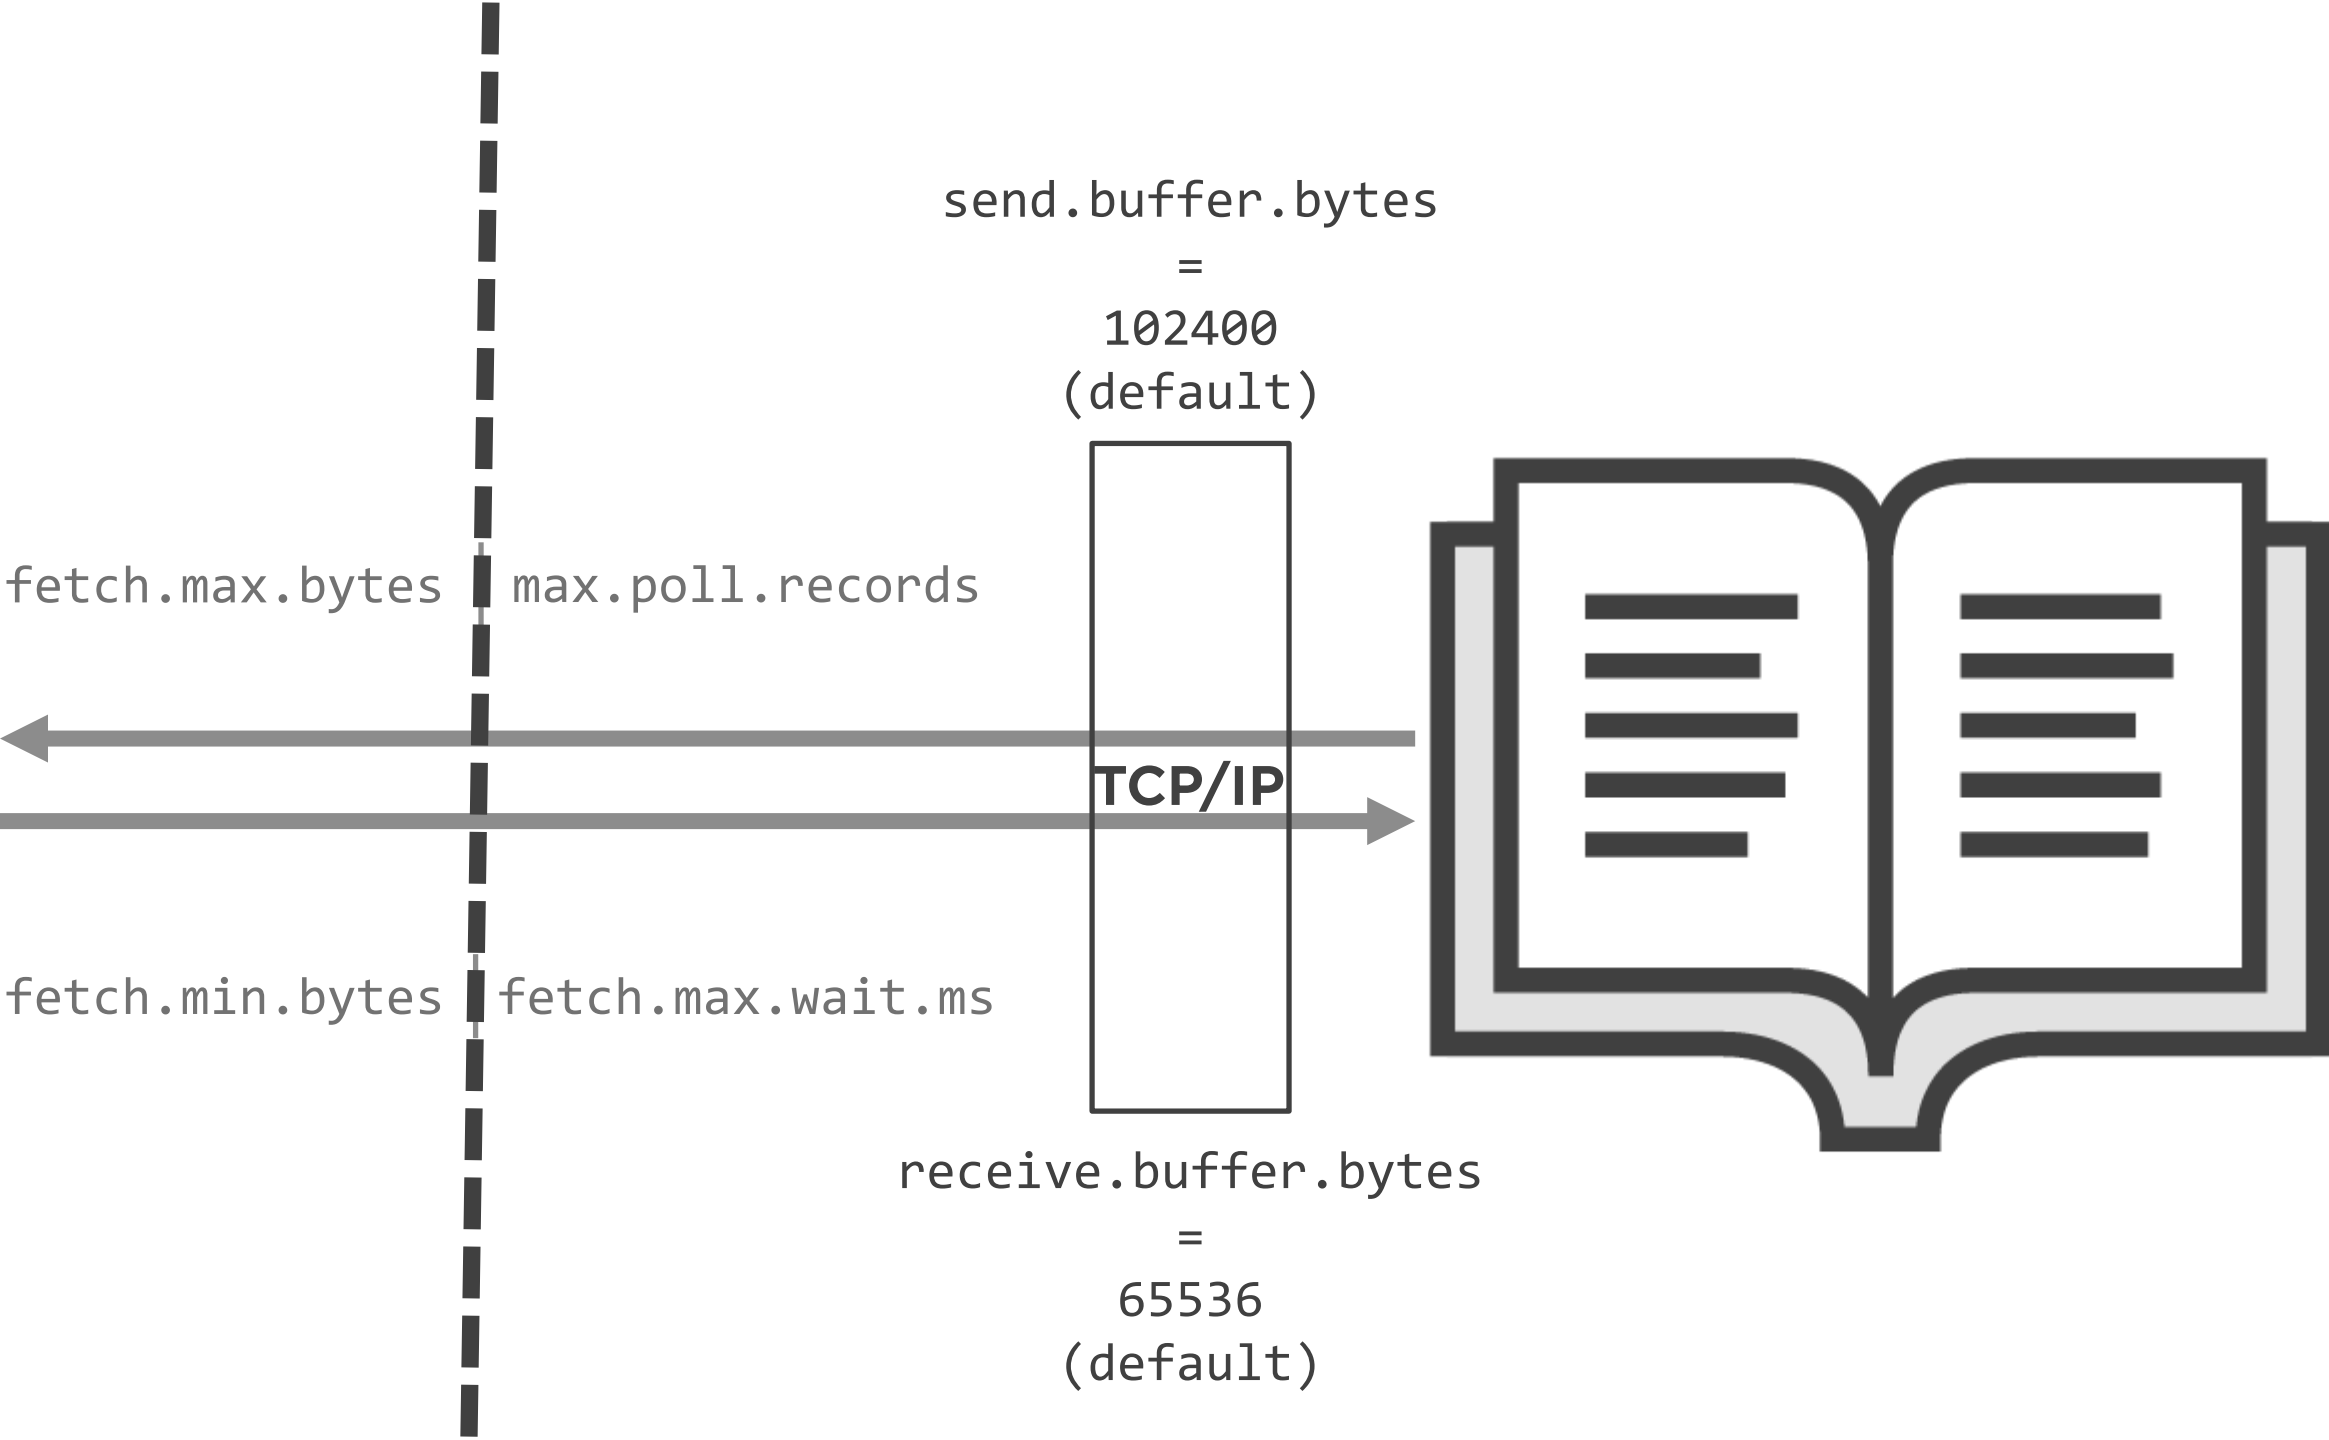
fetch.max.bytes | max.poll.records

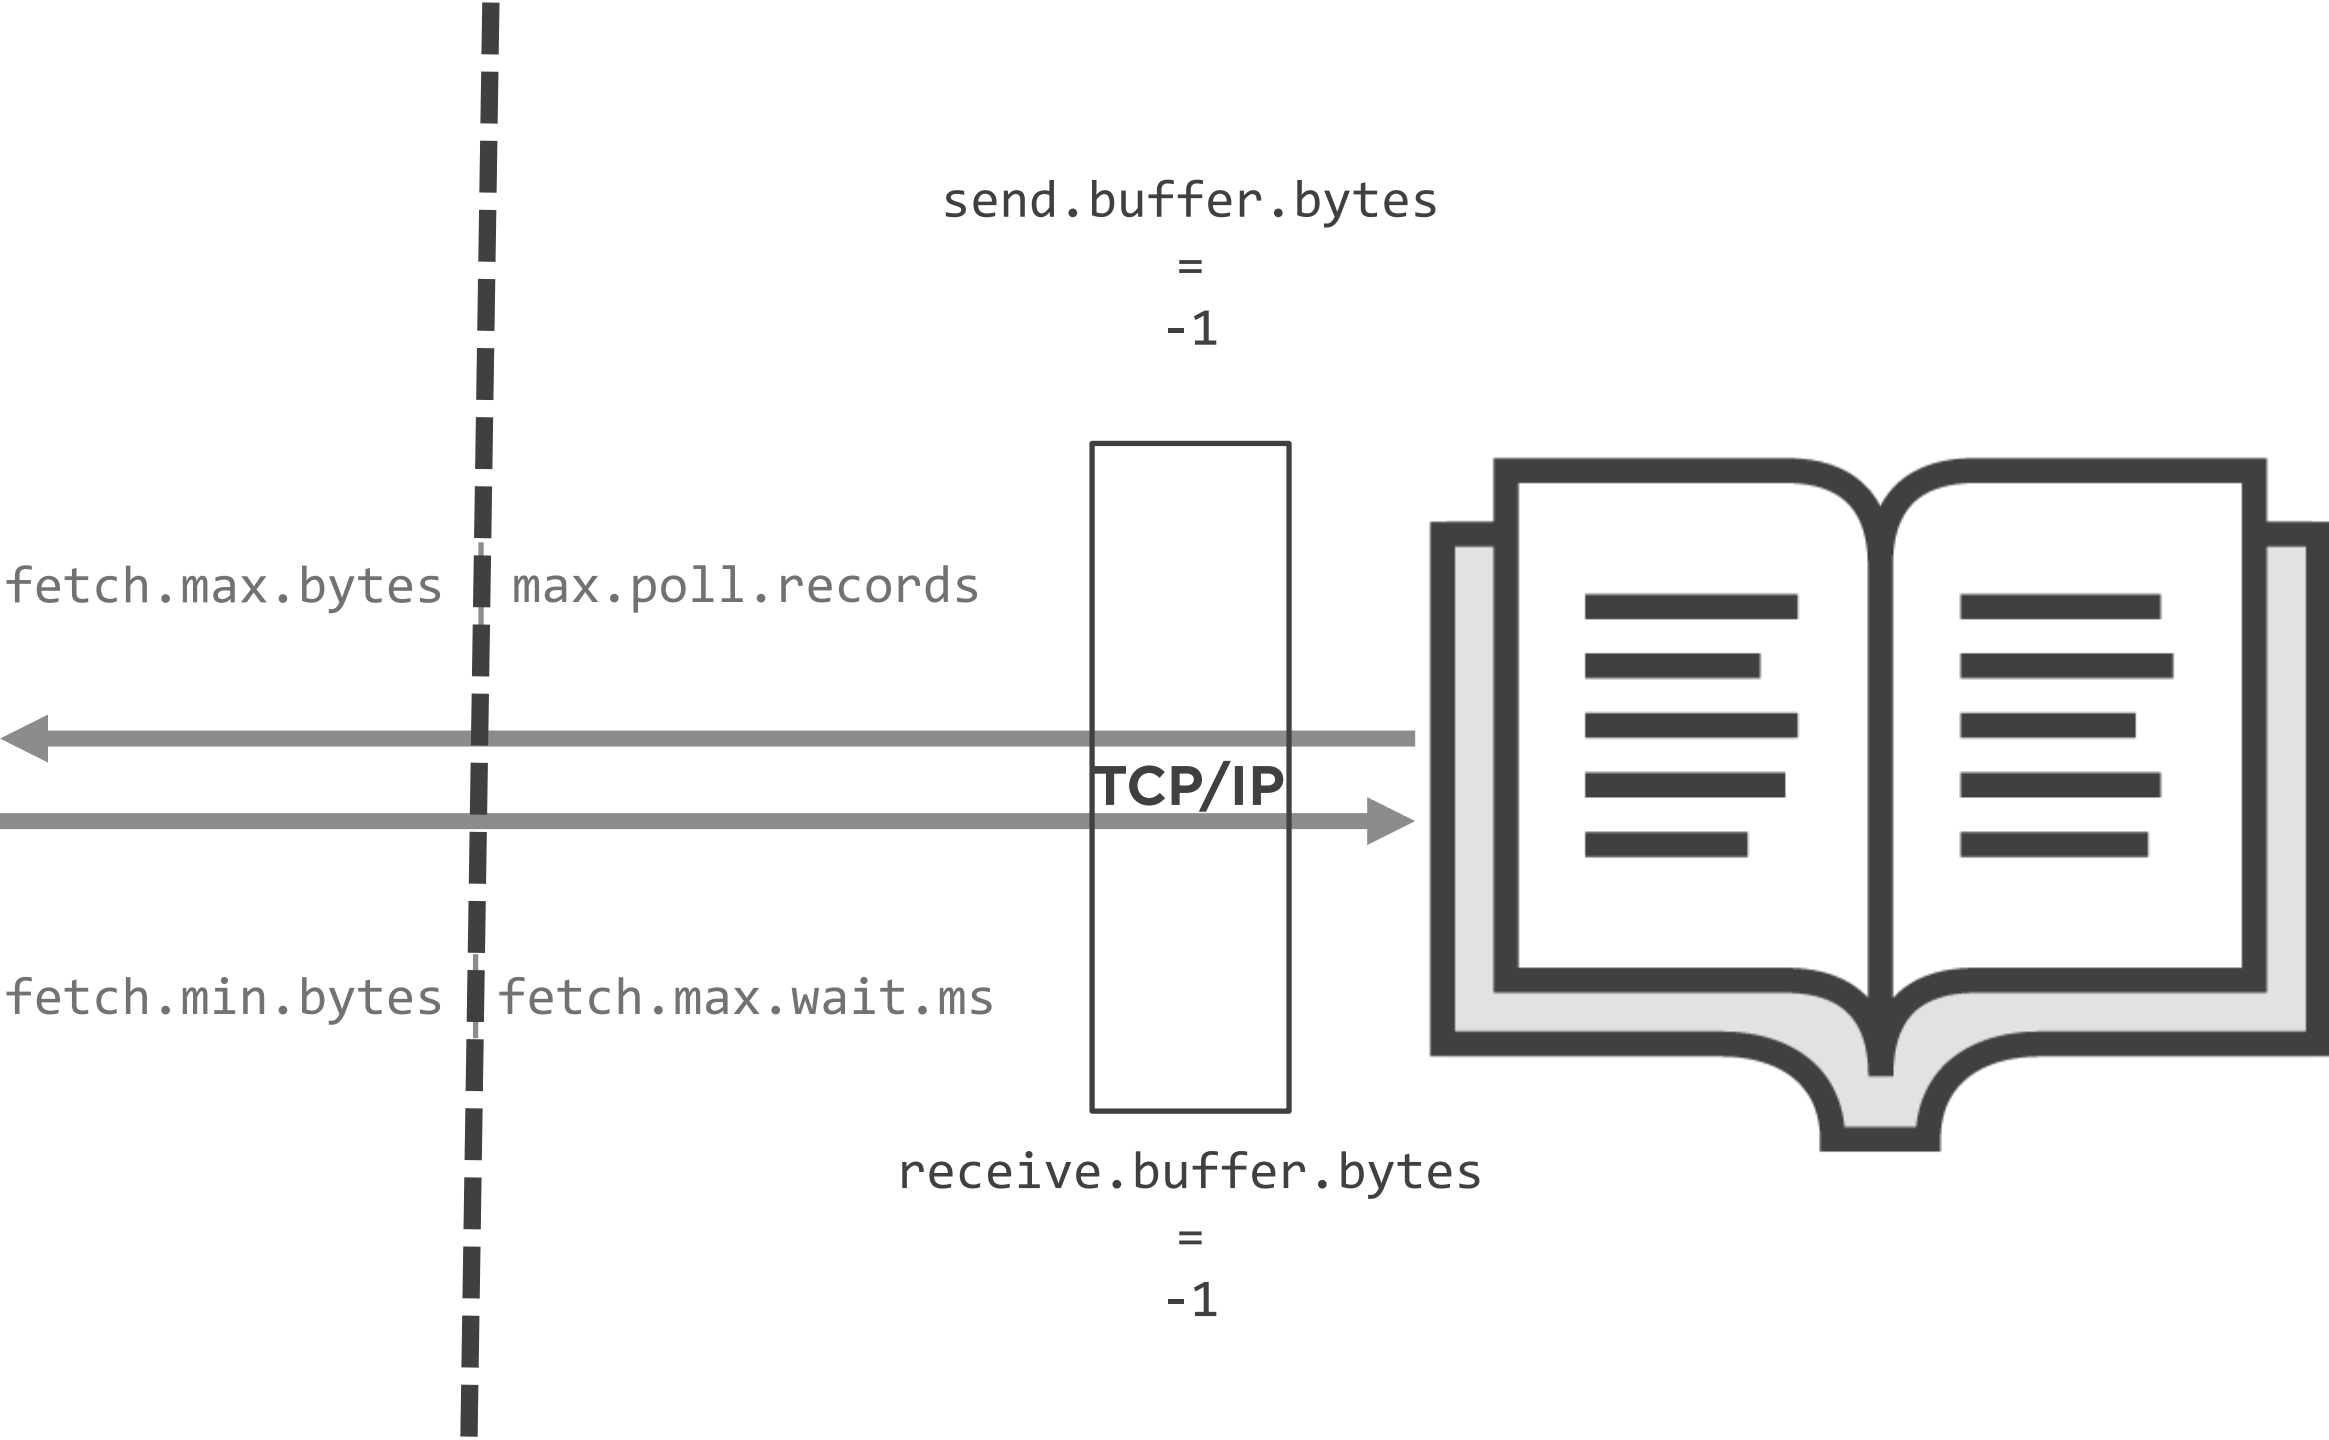


fetch.min.bytes | fetch.max.wait.ms

receive.buffer.bytes  
=  
65536  
(default)







`max.message.bytes` (topic)  
`message.max.bytes` (broker)

<

`max.partition.fetch.bytes`

=

1048576  
(default)


`send.buffer.bytes`

`fetch.max.bytes` | `max.poll.records`



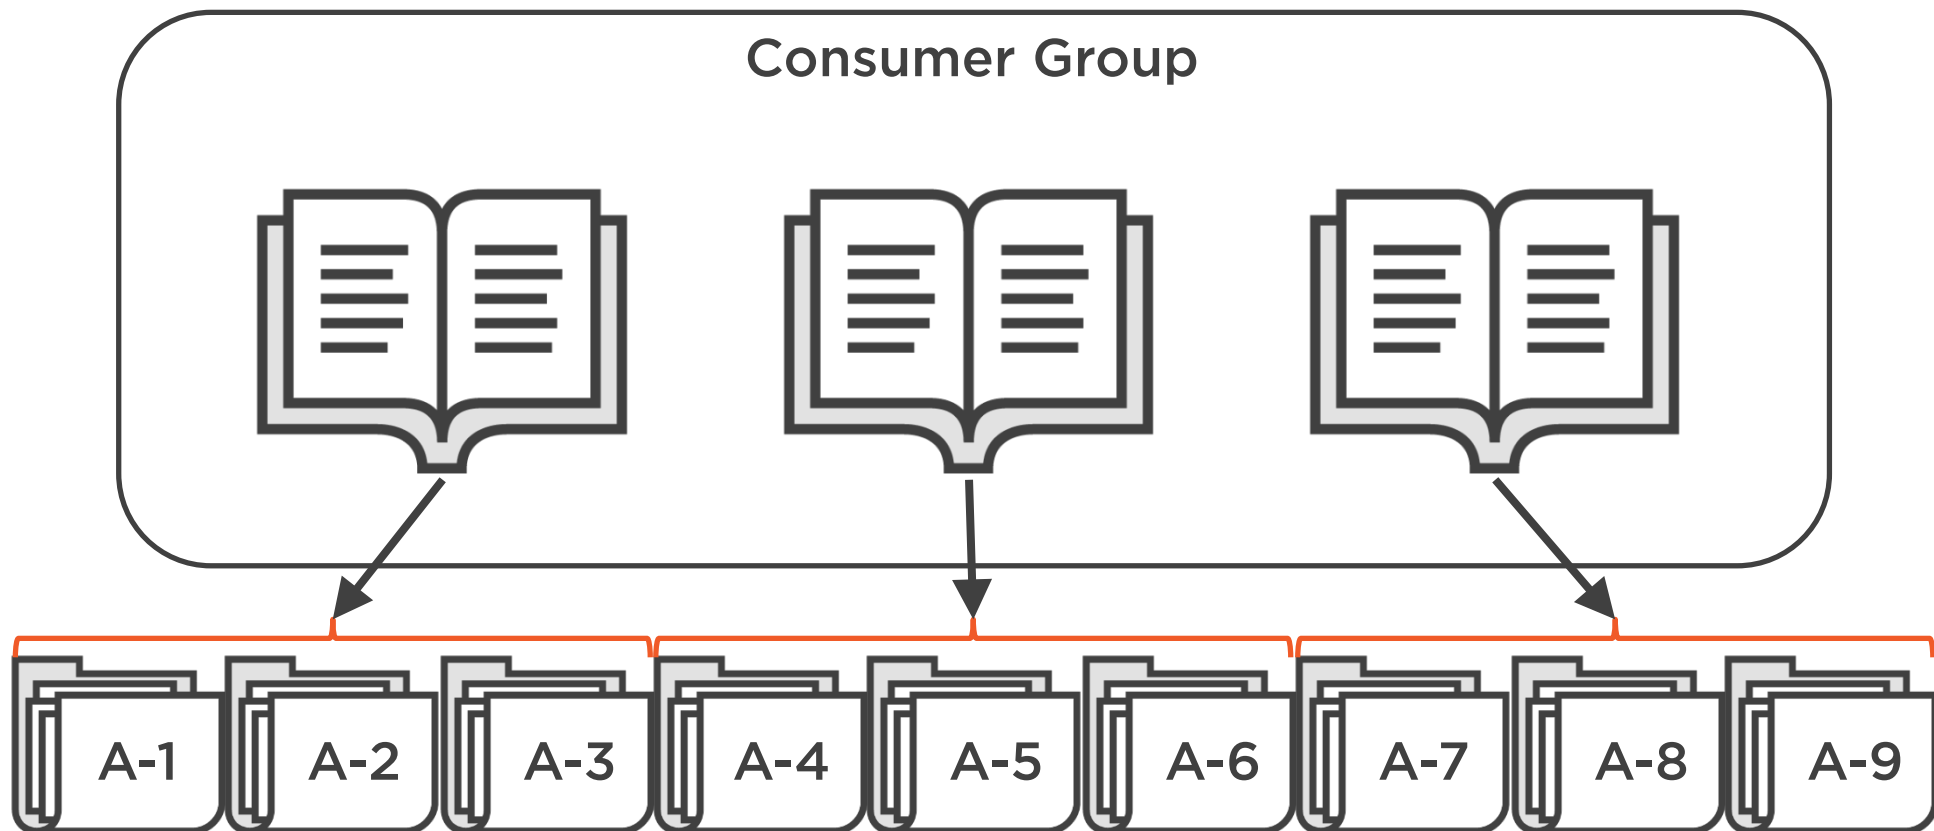
`fetch.min.bytes` | `fetch.max.wait.ms`

`receive.buffer.bytes`



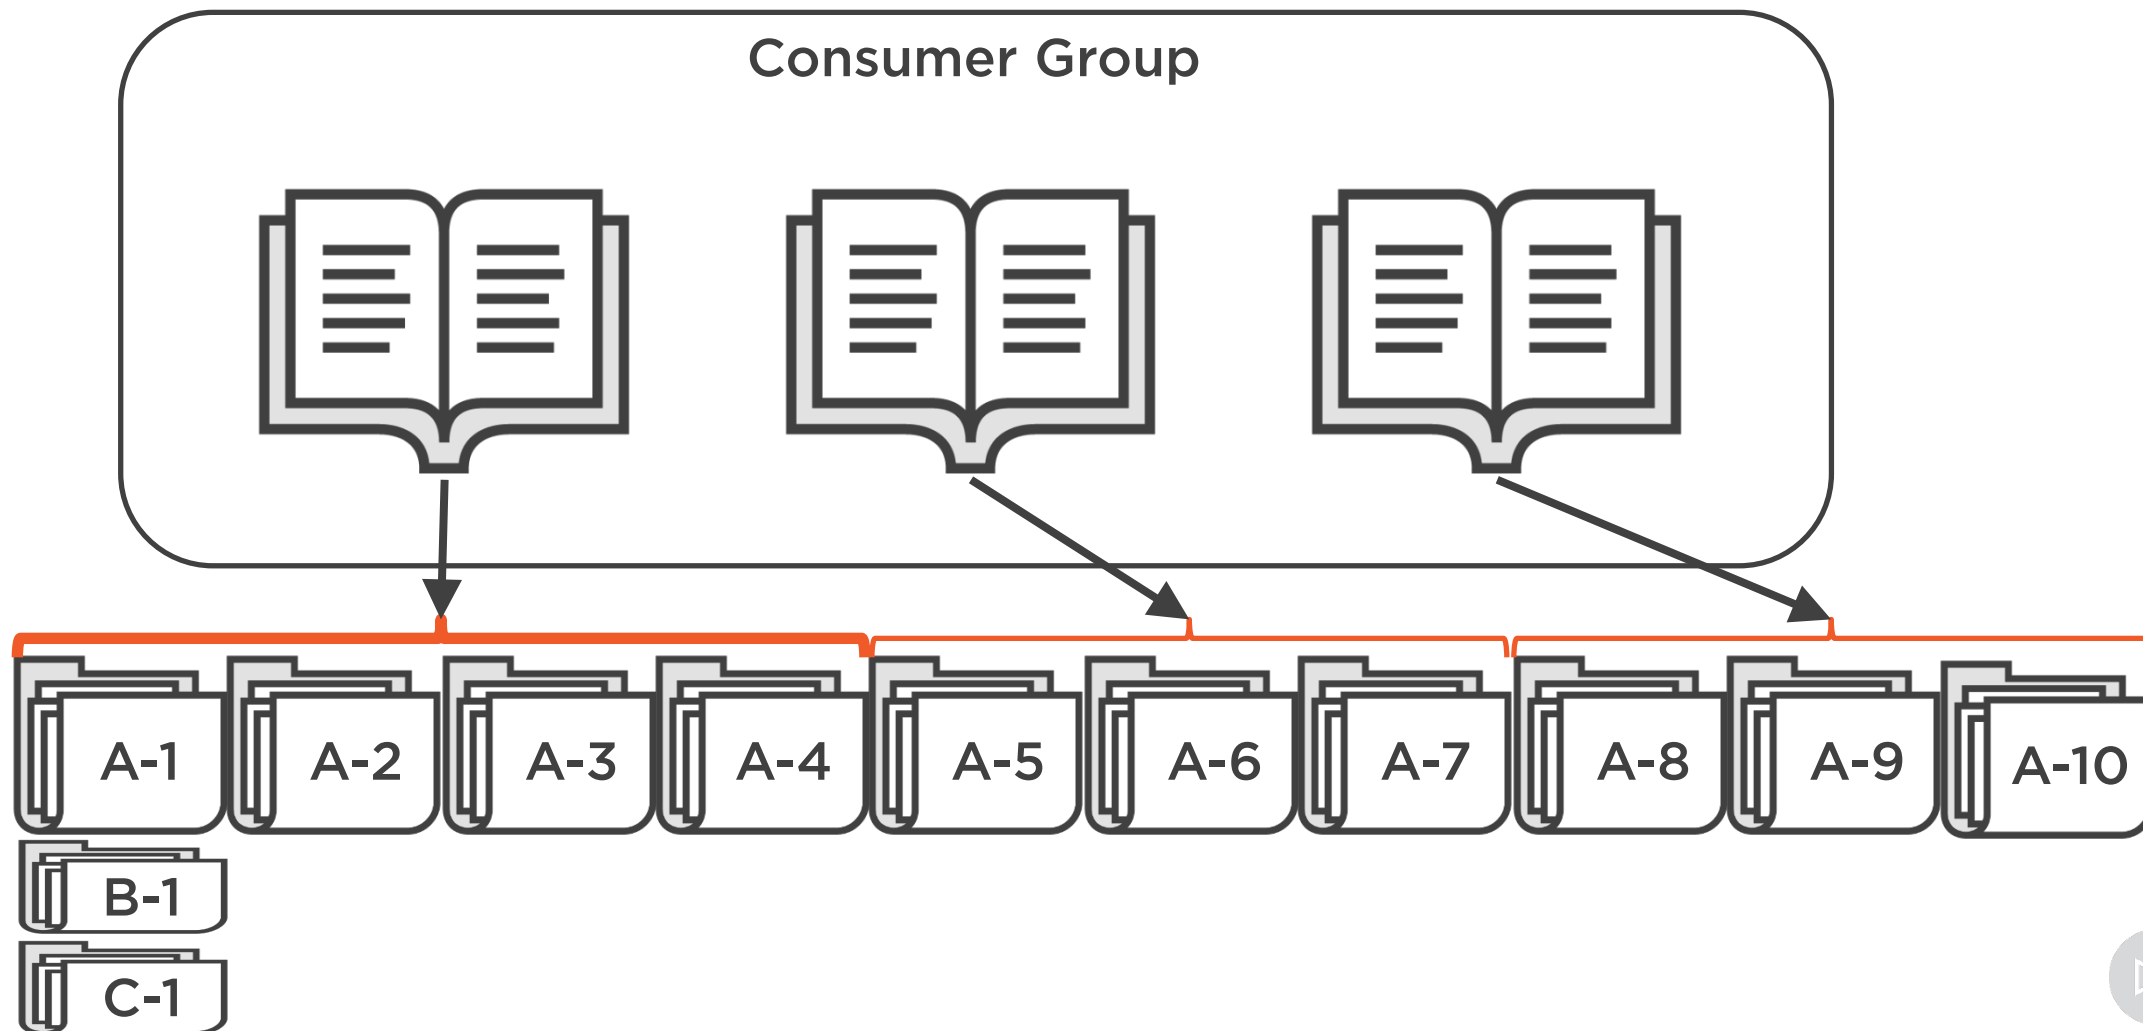
# partition.assignment.strategy

Default = `org.apache.kafka.clients.consumer.RangeAssignor`



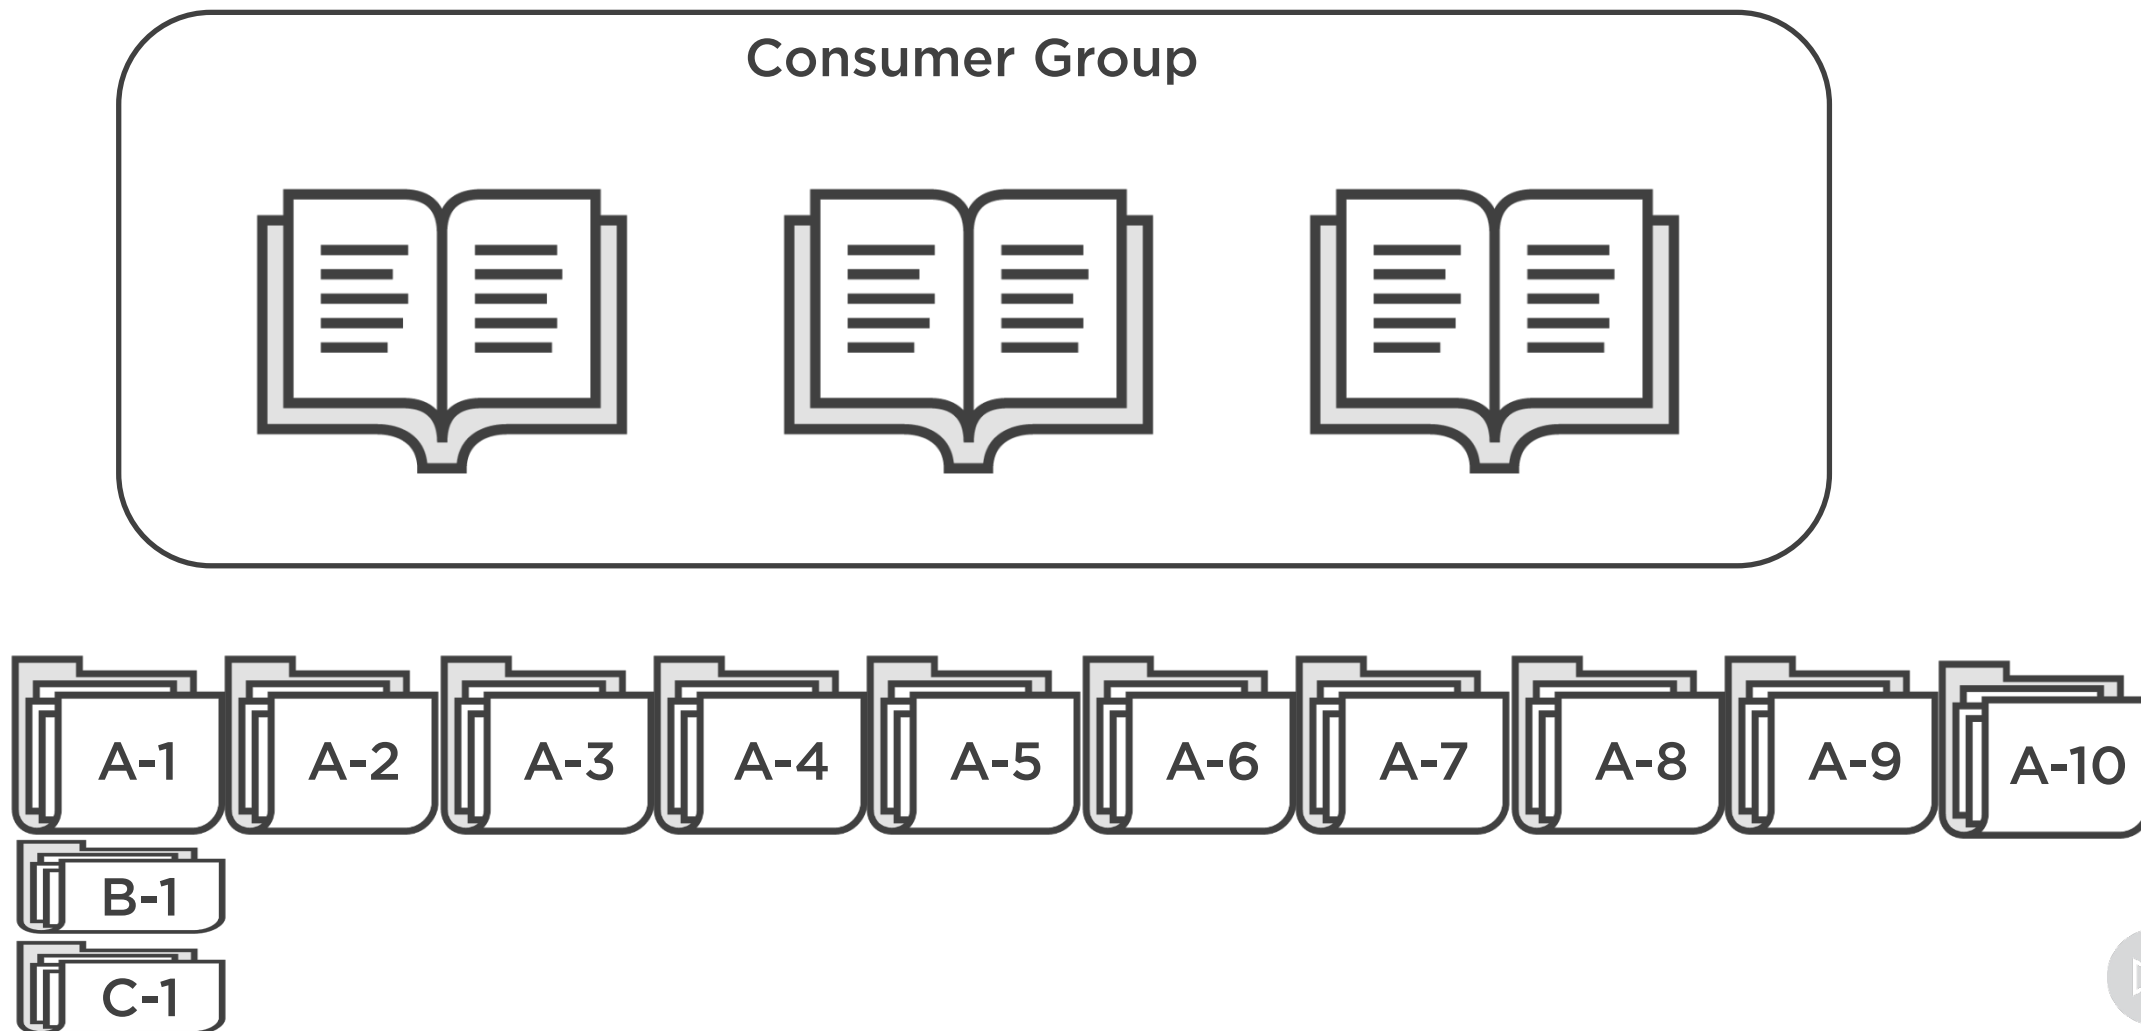
# partition.assignment.strategy

Default = `org.apache.kafka.clients.consumer.RangeAssignor`



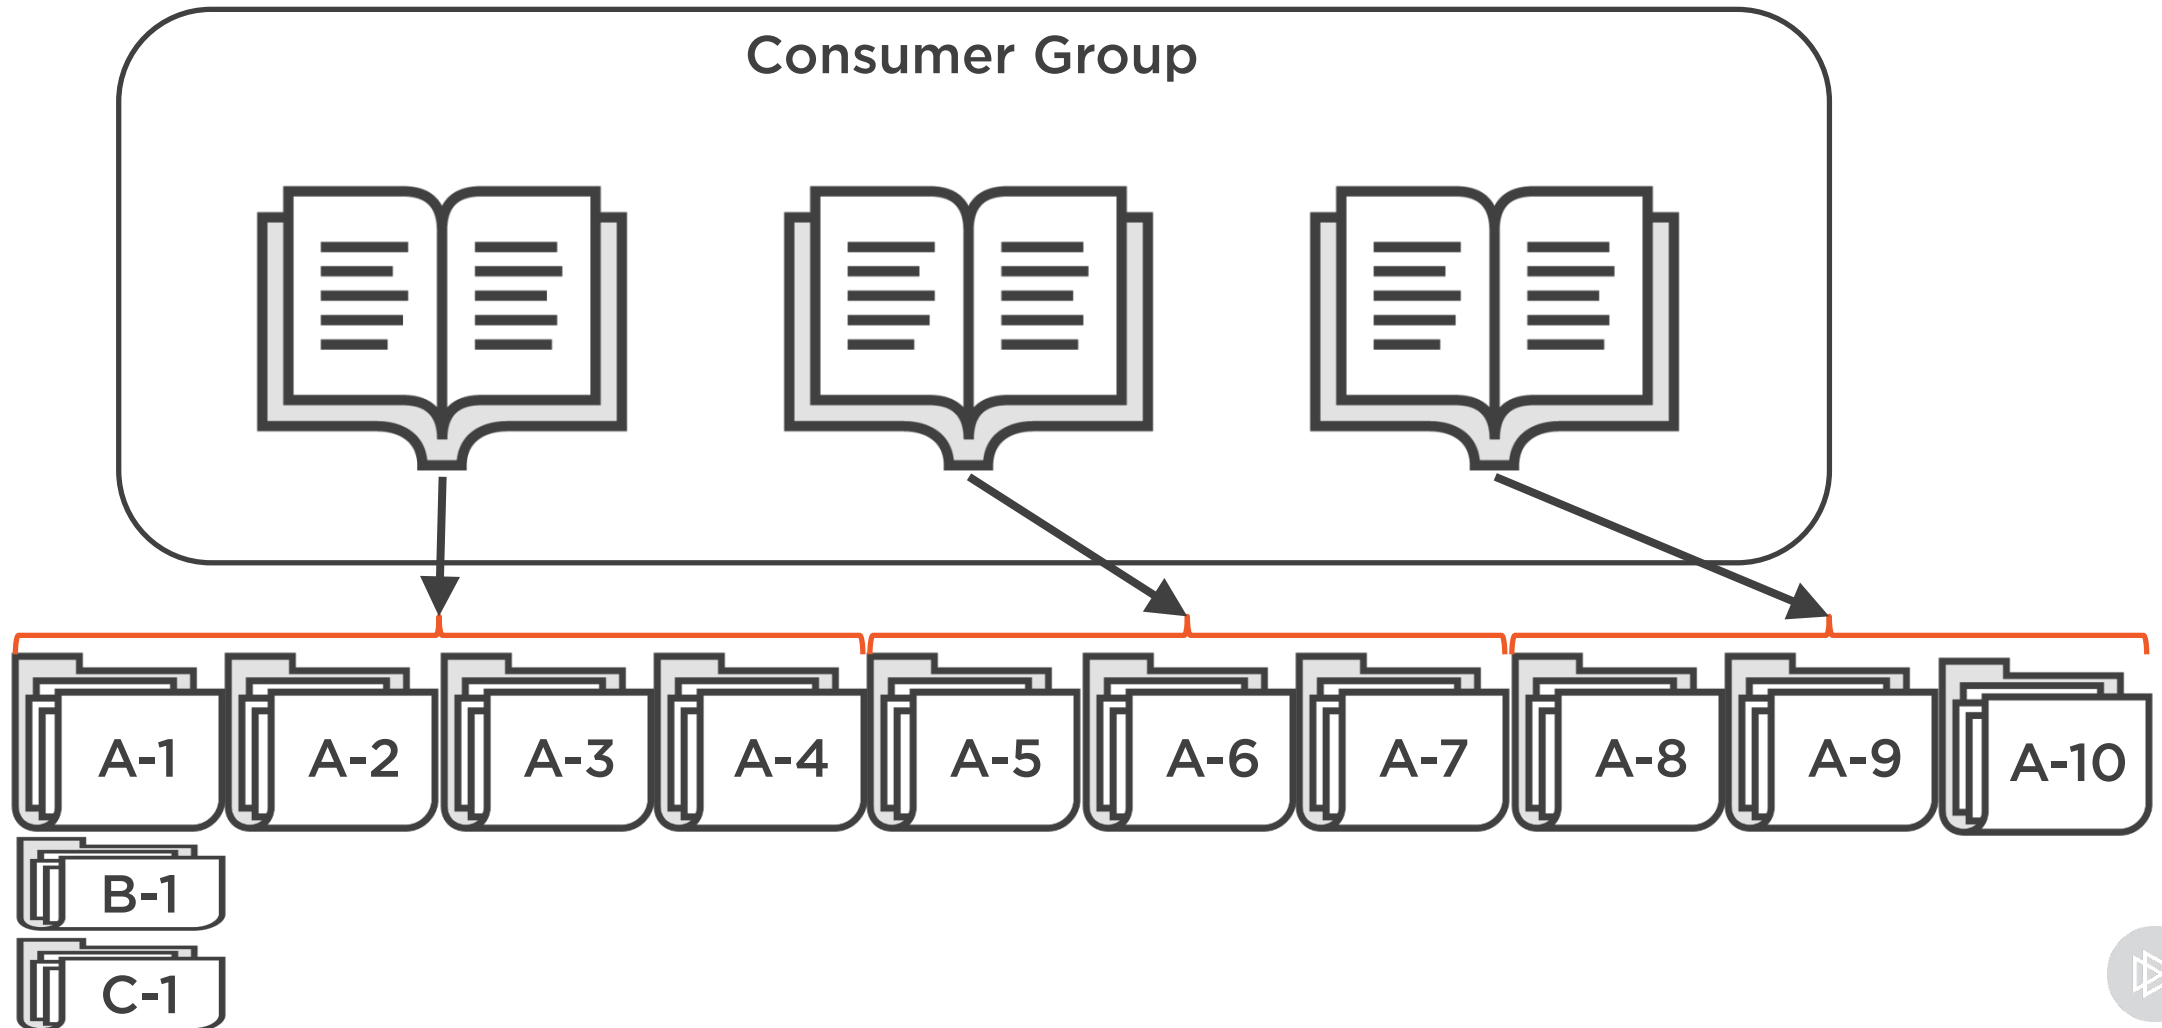
partition.assignment.strategy

org.apache.kafka.clients.consumer.RoundRobinAssignor



partition.assignment.strategy

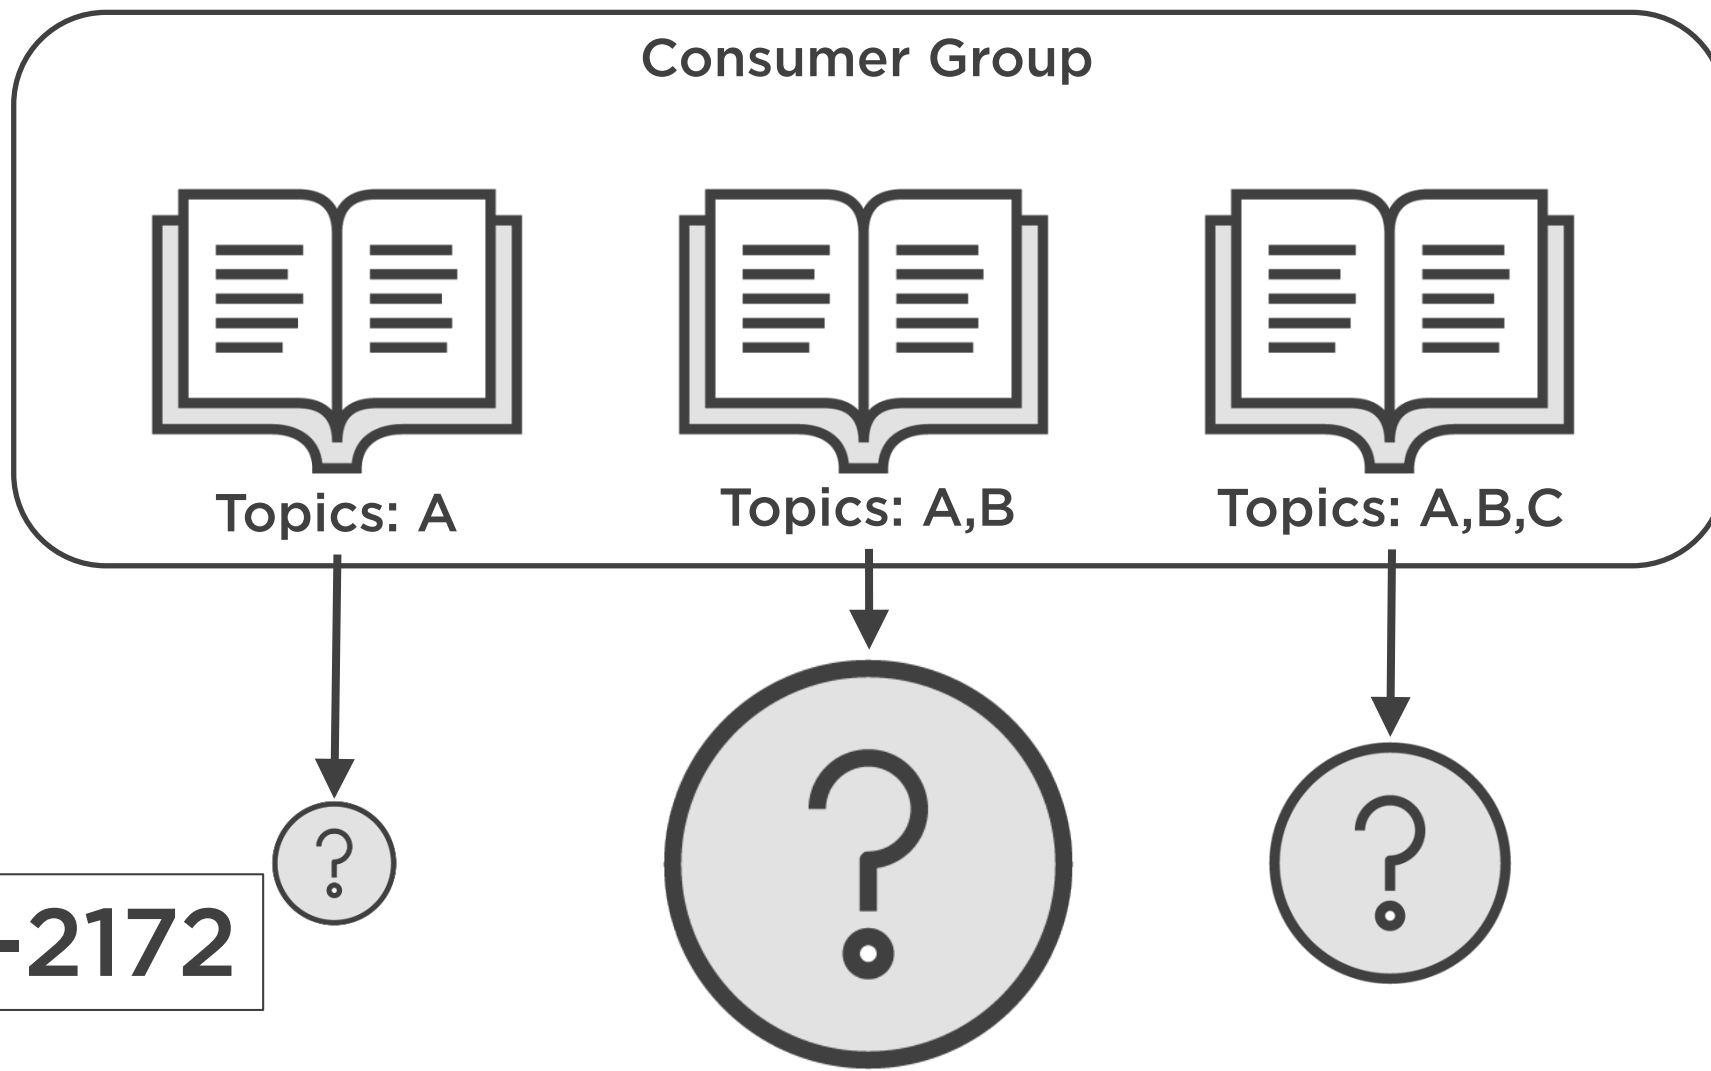
org.apache.kafka.clients.consumer.RoundRobinAssignor





partition.assignment.strategy

org.apache.kafka.clients.consumer.RoundRobinAssignor



KAFKA-2172




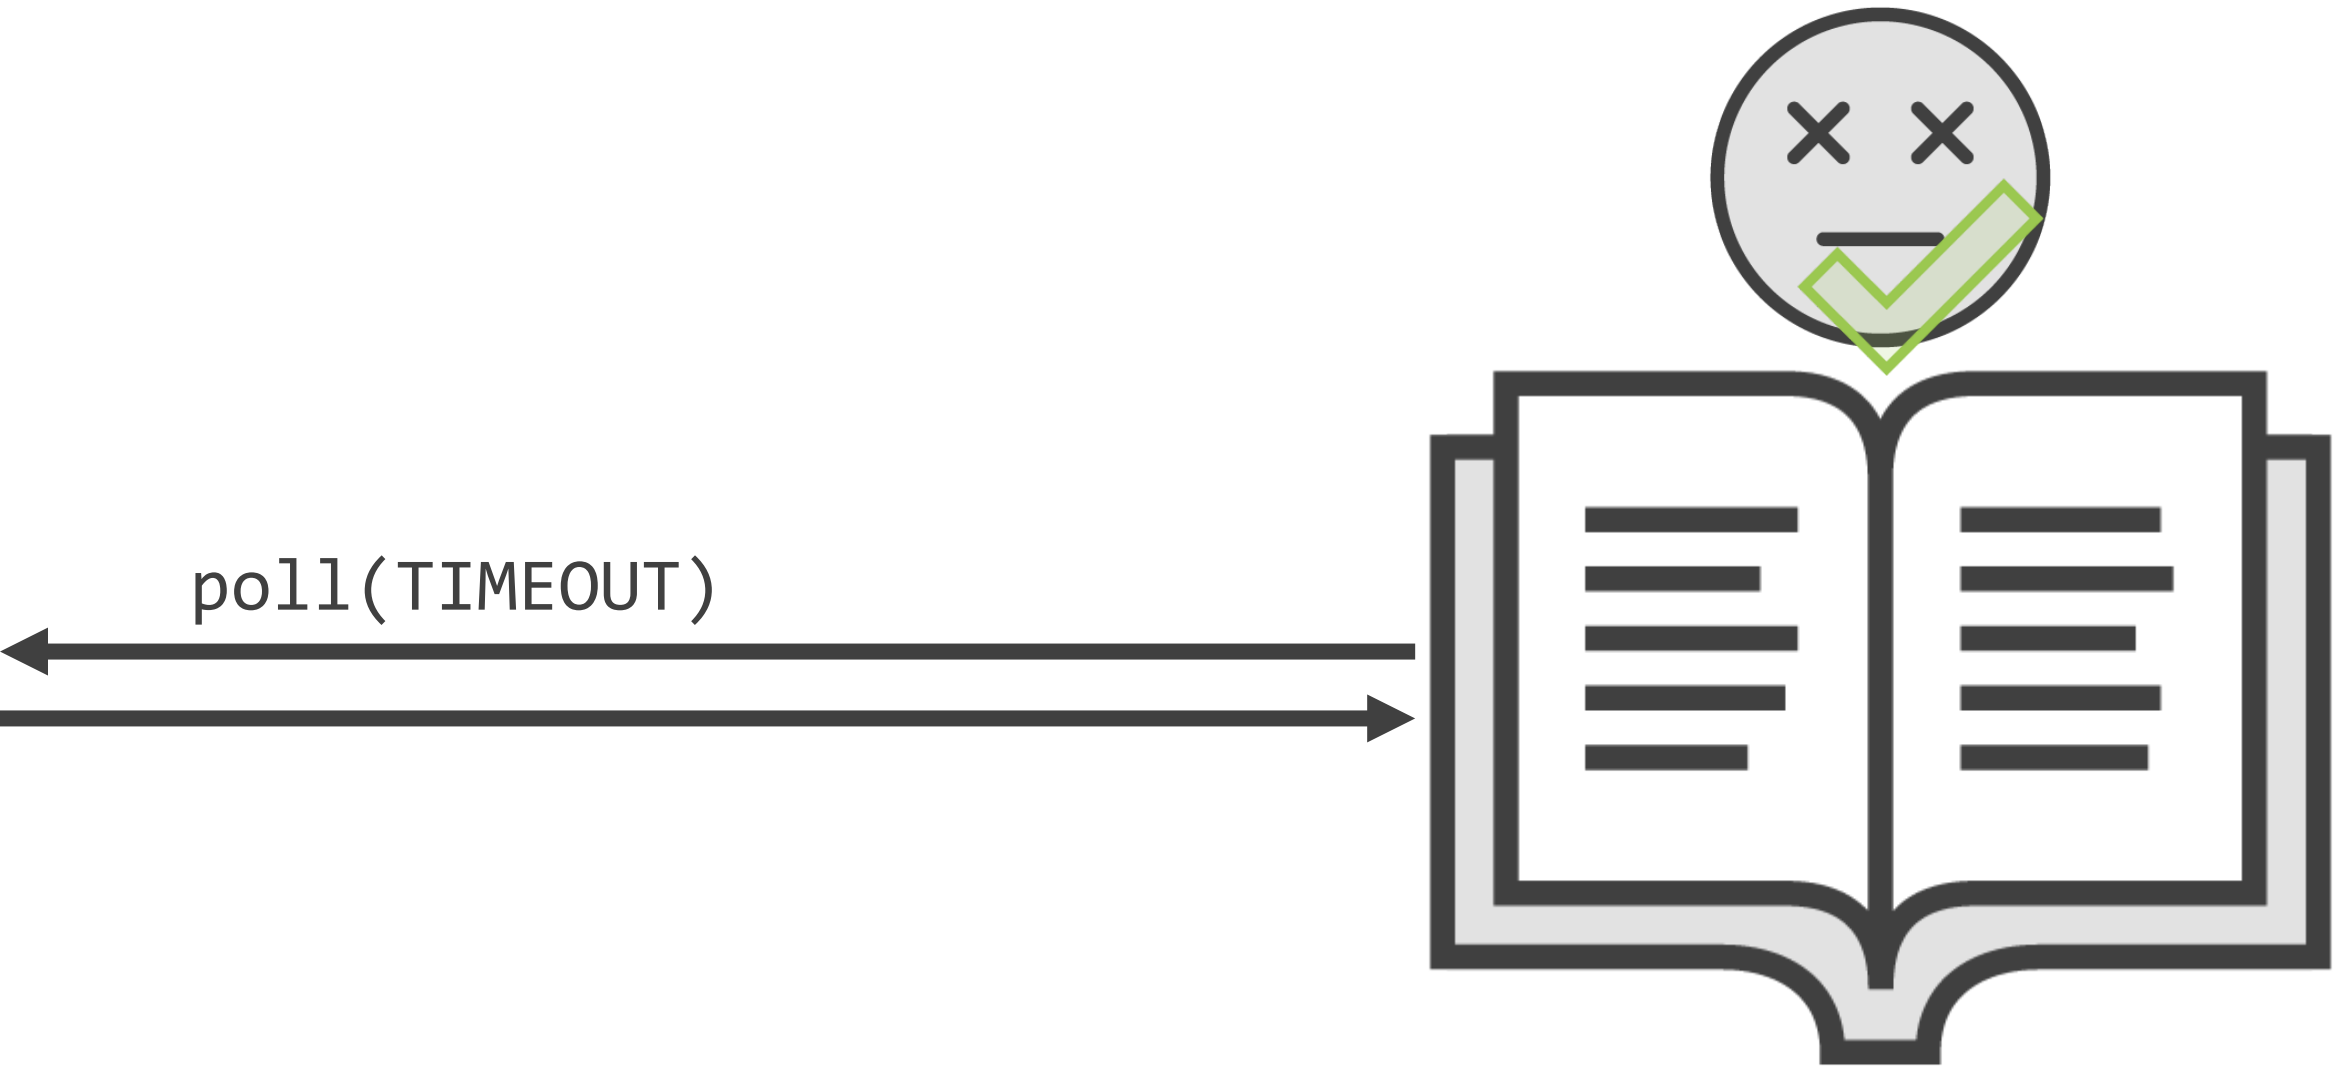
partition.assignment.strategy

Coming Soon? -> FairAssignor

org.apache.kafka.clients.consumer.internals.PartitionAssignor







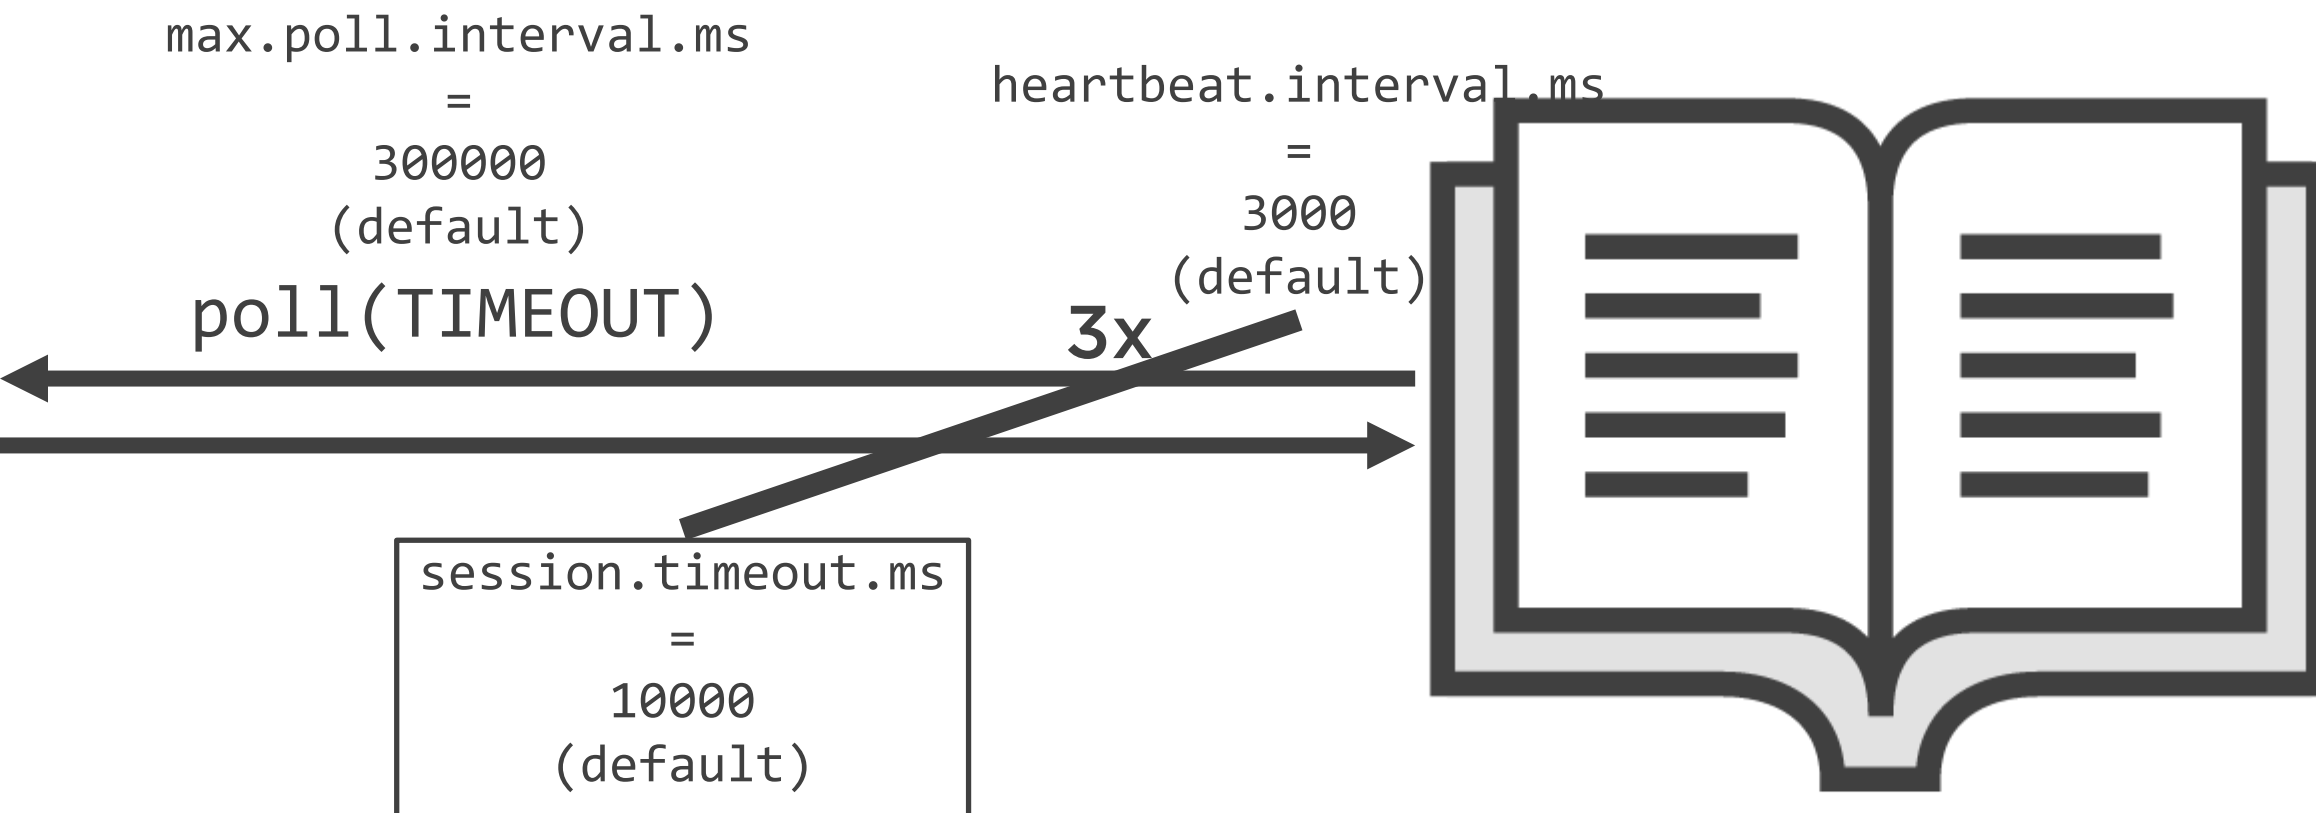
`max.poll.interval.ms`  
=  
300000  
(default)

`poll(TIMEOUT)`

`heartbeat.interval.ms`  
=  
3000  
(default)

`session.timeout.ms`  
=  
10000  
(default)






max.poll.interval.ms  
=  
300000  
(default)  
poll(TIMEOUT)

heartbeat.interval.ms  
=  
3000  
(default)

session.timeout.ms  
=  
10000  
(default)







`default.api.timeout.ms`

```
consumer.commitSync()  
    commitSync(offsets)  
    position(partition)  
    committed(partition)
```



`default.api.timeout.ms`

```
consumer.commitSync().....  
    commitSync(offsets)  
    position(partition)  
    committed(partition)
```



## KIP-266

`default.api.timeout.ms`

`=`

`60000`


`(default)`

`consumer.commitSync().....`

`commitSync(offsets)`

`position(partition)`

`committed(partition)`



# Configurations, Configurations, Configurations

- `[ssl/sasl].....`
- `client.dns.lookup`
- `request.timeout.ms`
- ...




# Consumer Offset Tracking


---



# Offset Tracking



# Offset Tracking



# Offset Tracking

`poll(TIMEOUT)`

Topic/Partition: A-1

**Position: 20**

`enable.auto.commit`

`=`

`true`  
`(default)`





# Offset Tracking

`poll(TIMEOUT)`

Topic/Partition: A-1

**Position: 50**

`enable.auto.commit`

`=`

`true`  
`(default)`



# Offset Tracking

```
enable.auto.commit  
=  
true  
(default)  
-----  
auto.commit.interval.ms  
=  
5000  
(default)
```



# Offset Tracking

**Overall Time: 4 sec**  
**Min Next Commit: 5 sec**

`poll(4 sec)`

`enable.auto.commit`

`=`

`true`

`(default)`

`auto.commit.interval.ms`

`=`

`5000`

`(default)`



# Offset Tracking

**Overall Time: 4 sec**  
**Min Next Commit: 5 sec**

poll(4 sec)

Topic/Partition: A-1  
**Position: 20**

`enable.auto.commit`

=

`true`

(default)

`auto.commit.interval.ms`

=

`5000`

(default)



# Offset Tracking

Overall Time: 8 sec  
Min Next Commit: 5 sec

```
enable.auto.commit  
    =  
    true  
    (default)  
-----  
auto.commit.interval.ms  
    =  
    5000  
    (default)
```



# Offset Tracking

**Overall Time: 8 sec**  
**Min Next Commit: 13 sec**

```
enable.auto.commit  
=  
true  
(default)  
-----  
auto.commit.interval.ms  
=  
5000  
(default)
```



# Offset Tracking

**Overall Time: 8 sec**  
**Min Next Commit: 13 sec**

`poll(4 sec)`

`enable.auto.commit`

`=`

`true`

`(default)`

`auto.commit.interval.ms`

`=`

`5000`

`(default)`



# Offset Tracking

Overall Time: **12 sec**  
Min Next Commit: **13 sec**

```
enable.auto.commit  
    =  
    true  
    (default)  
-----  
auto.commit.interval.ms  
    =  
    5000  
    (default)
```





# Offset Tracking

**Overall Time: 12 sec**  
**Min Next Commit: 13 sec**

poll(4 sec)

Topic/Partition: A-1  
**Position: 50**

`enable.auto.commit`

=

`true`

(default)

`auto.commit.interval.ms`

=

`5000`

(default)



# Offset Tracking

Overall Time: **16 sec**  
Min Next Commit: **13 sec**

```
enable.auto.commit  
    =  
    true  
    (default)  
-----  
auto.commit.interval.ms  
    =  
    5000  
    (default)
```



# Offset Tracking

`close()`

Topic/Partition: A-1  
**Position: 100**

`enable.auto.commit`

`=`

`true`  
`(default)`



# Consumer Offset Retrieval

---



`auto.offset.reset`

Topic/Partition: A-1  
Position: ???



`consumer.position(topicPartition)`



# auto.offset.reset

- latest (default)

Topic/Partition: A-1  
Position: ???



# auto.offset.reset

- latest (default)

Topic/Partition: A-1  
**Position: 51**



# auto.offset.reset

- latest (default)

Topic/Partition: A-1  
Position: ???





# auto.offset.reset

- latest (default)
- earliest

Topic/Partition: A-1  
**Position: ???**



# auto.offset.reset

- latest (default)
- earliest

Topic/Partition: A-1  
**Position: 5**

≠ 1



# auto.offset.reset

- latest (default)
- earliest

Topic/Partition: A-1  
**Position: ???**



# auto.offset.reset

- latest (default)
- earliest
- none

Topic/Partition: A-1  
**Position: ???**



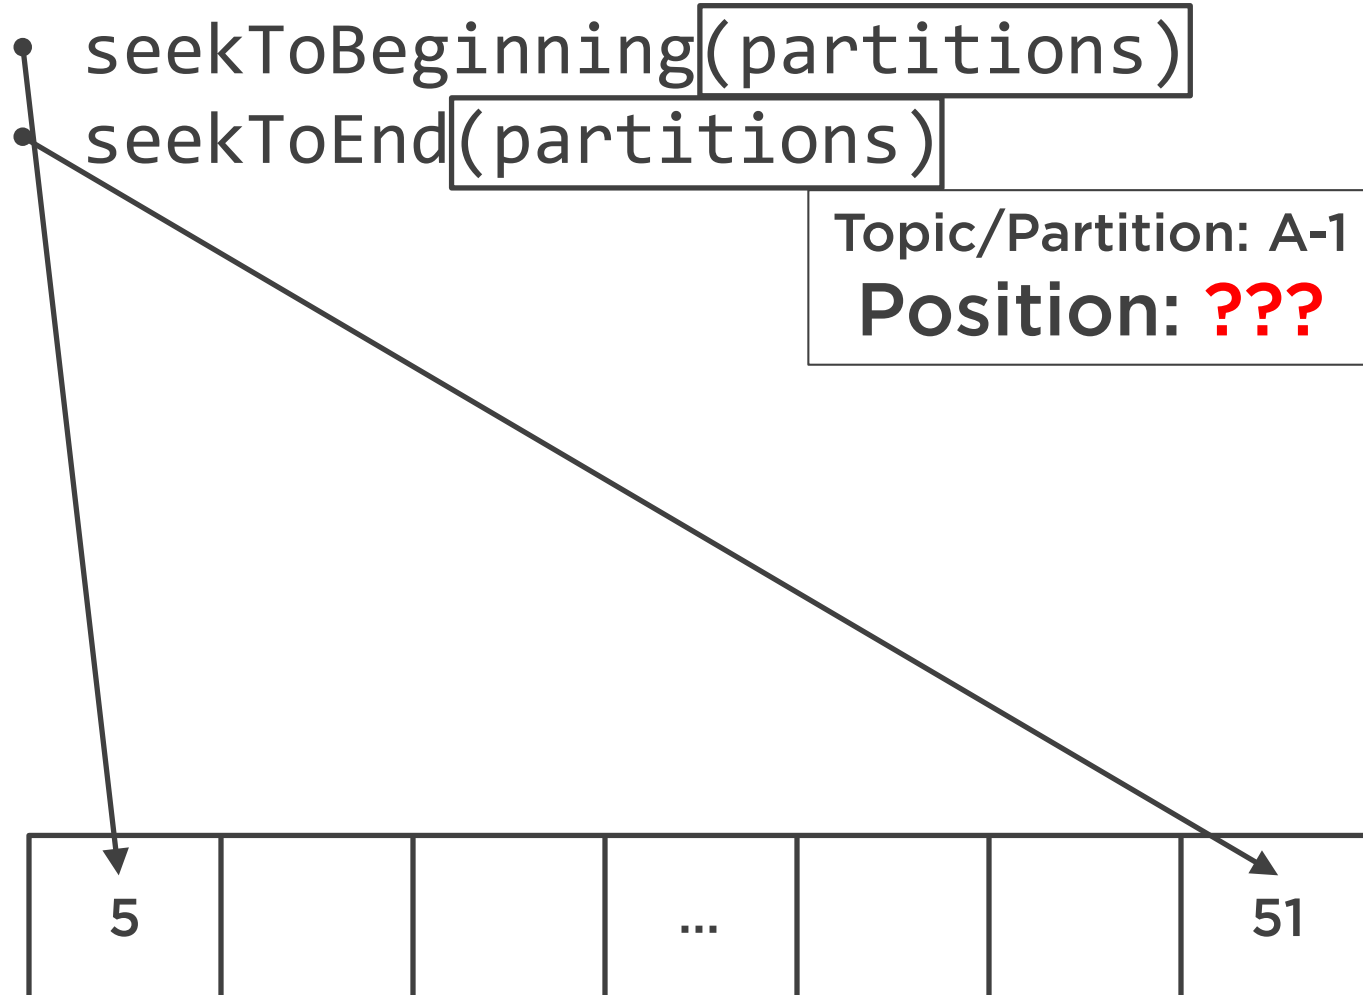
# auto.offset.reset

- latest (default)
- earliest
- none

Topic/Partition: A-1  
**Position:** 



# Offset Seeking



# Offset Seeking

- `seekToBeginning(partitions)`
- `seekToEnd(partitions)`

- `seek(partitions, position)`

Topic/Partition: A-1  
**Position: ???**

`poll(TIMEOUT)`



# Offset Seeking

- `seekToBeginning(partitions)`
- `seekToEnd(partitions)`

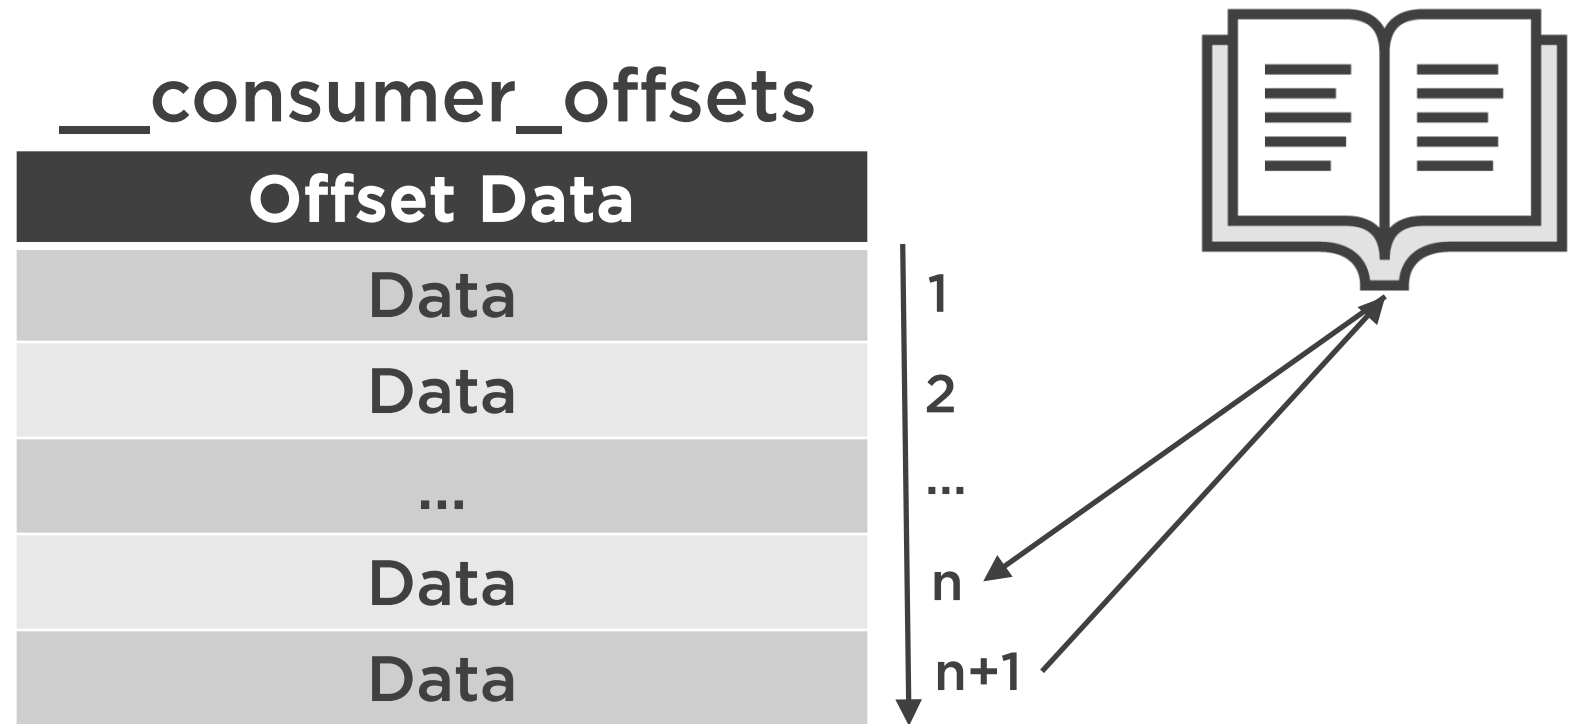
- `seek(partitions, position)`

Topic/Partition: A-1  
**Position: 9**

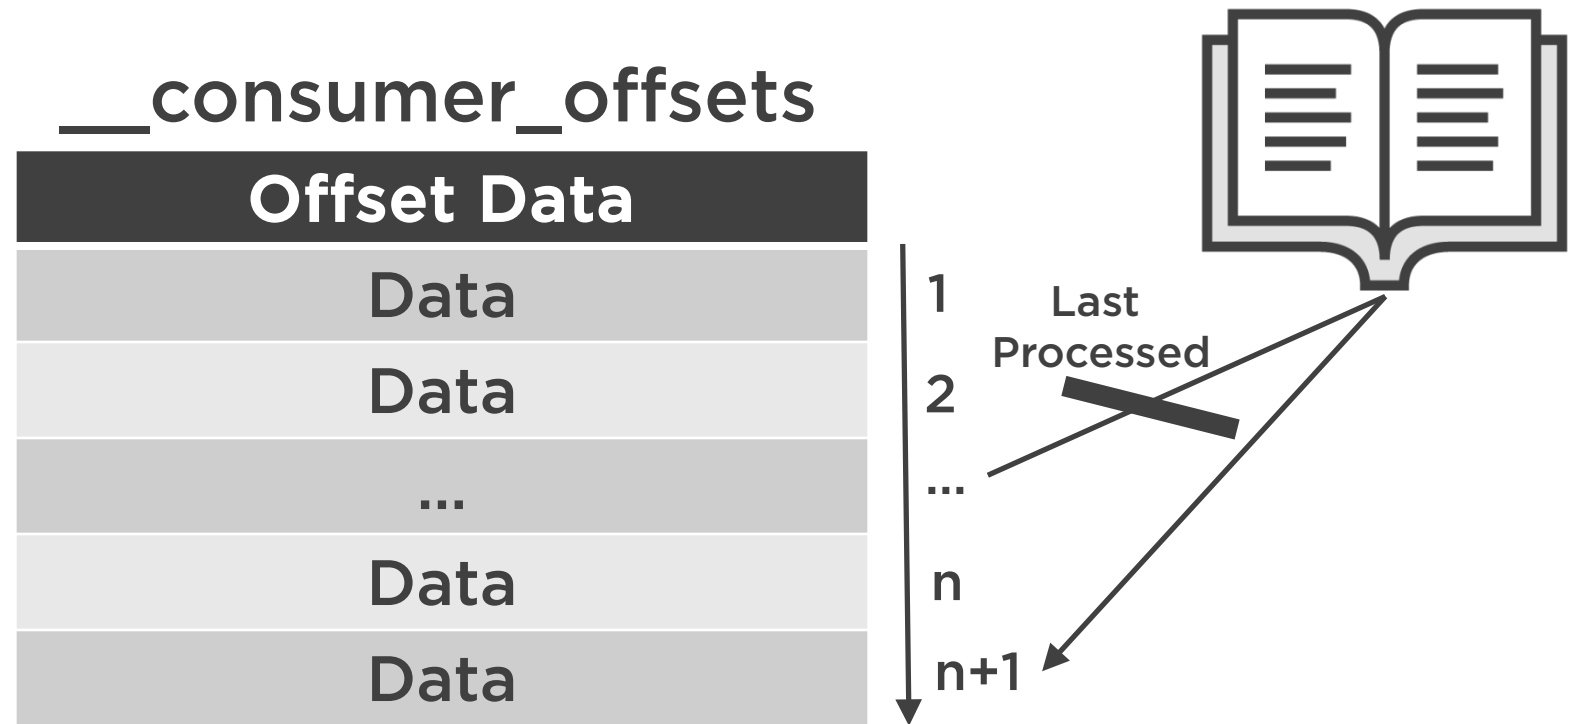




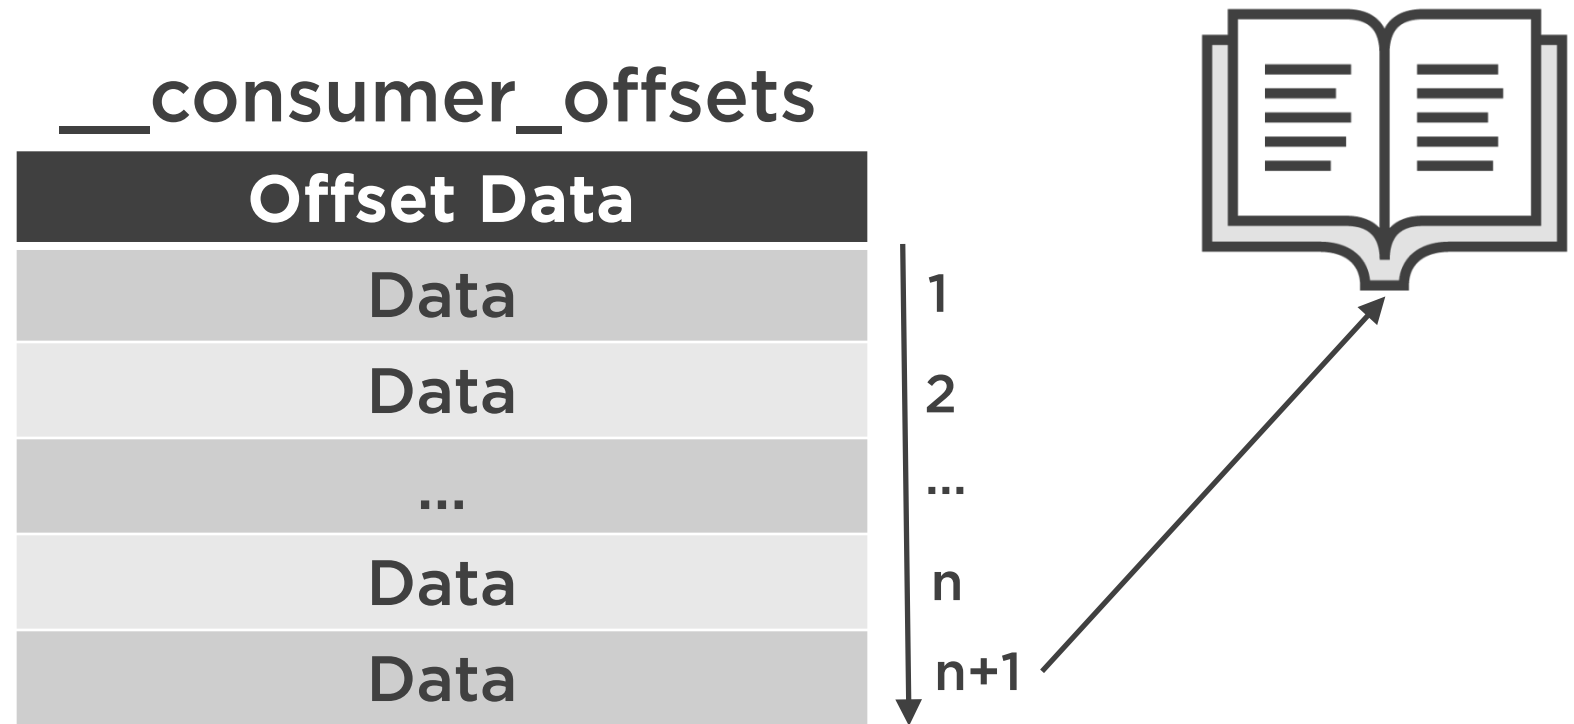
# Offset Tracking



# Offset Tracking



# Offset Tracking





# Reliability in Production


---




# Reliability







acks = all (or -1)





acks = all (or -1)


$\approx 1$  ???

min.insync.replicas


=

1

(default)



## Cluster




acks = all (or -1)

`min.insync.replicas`  
=  
2


1 ISR < 2 Min ISR



NotEnoughReplicas..Exception




## Cluster



acks = all (or -1)



min.insync.replicas  
=  
2

2 ISR = 2 Min ISR



acks = all (or -1)


min.insync.replicas  
=  
2



`acks = all (or -1)`

`min.insync.replicas`  
=  
2

broker = topic default  
OR  
topic override



## Cluster



# Cluster



# Cluster



# Cluster







# Cluster



# Cluster







# Offsets

enable.auto.commit  
auto.commit.interval.ms



consumer.commitSync(offsets)


versus

consumer.commitSync(|)

alloffsets




# Offsets




# auto.offset.reset

- latest
- earliest
- none



# auto.offset.reset

- latest
  - earliest
  - none
- 
- seek(partitions, position)
- ???
- 






# Reliability Validation

- Reliable Configurations
  - `org.apache.kafka.tools.VerifiableProducer`
  - `org.apache.kafka.tools.VerifiableConsumer`
- Test Scenarios
  - Consumer Rebalancing
  - Leader re-election
  - [Consumer/Producer/Broker] rolling restart
- Monitoring!





compression.type  
=  
uncompressed  
(default)






- snappy
- lz4
- gzip
- zstd

# Empowerment through Custom Messages: Serializers & Deserializers

---








0	1	1	1	0	1	1	1	1	0
0	0	1	1		0	0	0	1	1
1	0	1	0	1	1	1	0	1	1
1	0	0	1	0	1	1	1	0	1
0	1	0	1	0	0	1	1		0
0	1	1	0	1	0	1	0	1	1









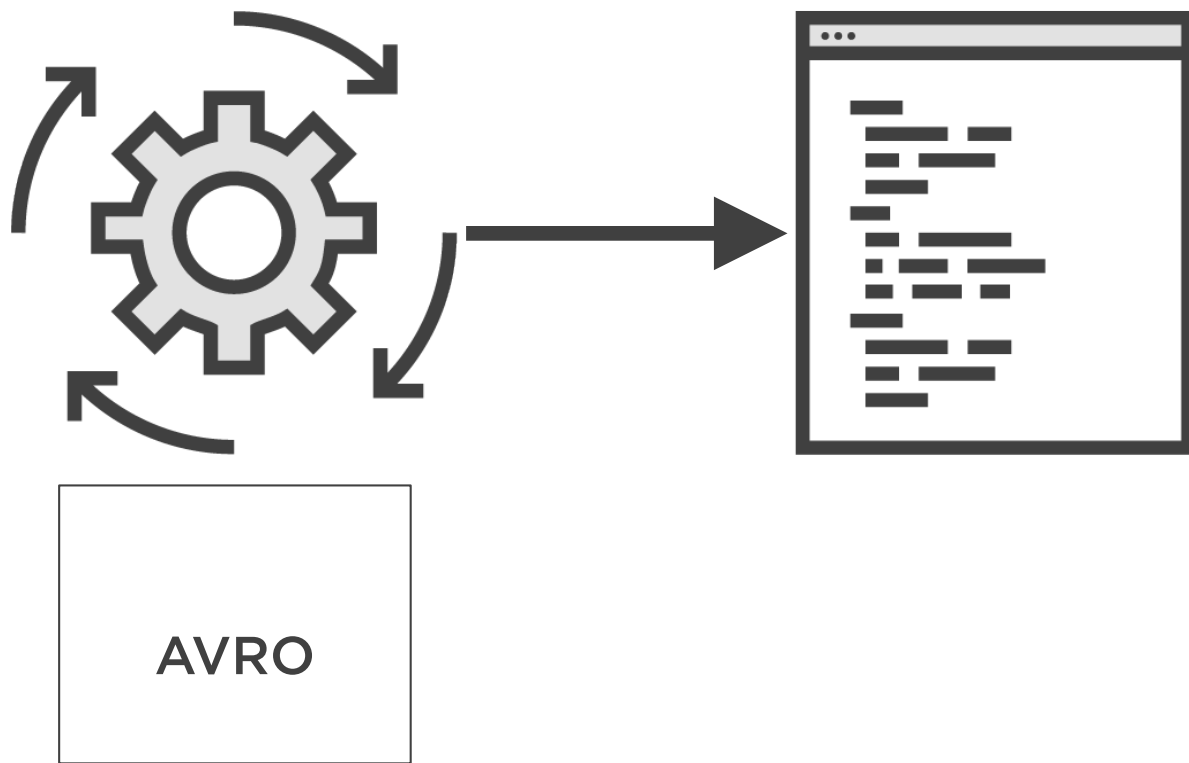





# Cluster













Version	Schema
1	A
2	A
1	B
2	B
3	B





# Resources

- **Configuring Apache Kafka for Performance and Resource Management**
  - [cloudera.com/documentation/kafka/latest/topics/kafka\\_performance.html](https://cloudera.com/documentation/kafka/latest/topics/kafka_performance.html)
- **Running Kafka in Production**
  - [docs.confluent.io/current/kafka/deployment.html](https://docs.confluent.io/current/kafka/deployment.html)
- **Decoupling Systems with Apache Kafka, Schema Registry and Avro**
  - [confluent.io/blog/decoupling-systems-with-apache-kafka-schema-registry-and-avro](https://confluent.io/blog/decoupling-systems-with-apache-kafka-schema-registry-and-avro)
- **Avro Documentation**
  - [avro.apache.org/docs/current](https://avro.apache.org/docs/current)
- **Schema Registry Installation Documentation**
  - [docs.confluent.io/current/schema-registry/docs/installation.html#development](https://docs.confluent.io/current/schema-registry/docs/installation.html#development)



# Summary



## Configuration

- Producer
- Consumer

## Reliability

## Message Typing via Avro

